

Project 2 - Emotion

Henrik Daniel Christensen^[hench13@student.sdu.dk]
Frode Engtoft Johansen^[fjoha21@student.sdu.dk]

DM873: Deep Learning
University of Southern Denmark, SDU
Department of Mathematics and Computer Science

1 Introduction

The aim of this project is to design and implement two deep learning models for text/sentiment classification, capable of identifying the emotion conveyed in a given text. The model is trained on the *dair-ai/emotion* dataset, a widely used collection of 20,000 labeled tweets available through Hugging Face. This dataset is divided into six emotion categories: sadness (0), joy (1), love (2), anger (3), fear (4), and surprise (5). Of the total dataset, 2,000 tweets are reserved exclusively for final testing.

2 Label Distribution

We started by analyzing the distribution of the labels in the dataset, see Figure 1.




Fig. 1. Distribution of the labels in the dataset.

From the figure, it is evident that the dataset is highly imbalanced, with the majority of tweets labeled as either joy or sadness. However, the validation set exhibits a roughly similar distribution of labels, suggesting that the test set is likely to follow a comparable pattern.

This imbalance in the training data is expected to impact the model's performance, particularly for less represented labels, as sentences labeled with joy are less likely to be classified correctly. This challenge must be carefully considered when evaluating the model's performance and interpreting the results.

3 Preprocessing

To enable a machine learning model to process textual data effectively, it is essential to convert the text into numerical representations. The first step in this process is tokenizing the text, which involves splitting sentences into smaller units, such as words. For this task, we utilize a custom `RegexTokenizer` from the NLTK library, which filters the text to retain only words, numbers, and a limited set of special characters. This ensures that irrelevant symbols are excluded while preserving the meaningful components of the text.

Following tokenization, the text is further refined by removing stopwords and applying stemming. Stopwords, which are frequently used words such as "the," "is," and "and," are removed as they are unlikely to contribute significant value to sentiment analysis. This is particularly important when using a limited sequence length for input sentences, as it allows the model to focus on more informative features while reducing the size of the vocabulary. Consequently, the model learns fewer parameters, which can improve its generalization and efficiency.

In addition, we apply stemming using the PorterStemmer from the NLTK library, which reduces words to their base or root forms (e.g., "feeling" becomes "feel"). This process further reduces the vocabulary size by grouping inflected or derived word forms together, helping the model to recognize similar meanings and improve its performance.

To better understand the text data after preprocessing, we visualize the most common words using a WordCloud, as shown in Figure 2. This provides an intuitive way to analyze the prominent terms in the dataset and assess whether the preprocessing aligns with the goals of the sentiment analysis task.




Fig. 2. WordCloud of the vocabulary.

As expected many emotional words are present in the vocabulary, e.g. 'love', 'friend', 'hate', 'never', 'go', etc.

To be able to process the text data, we need to choose a maximum sentence length. A boxplot of the distribution of the lengths of the sequences can be seen in Figure 3.




Fig. 3. Distribution of the lengths of the sequences.

From the boxplot, we see that the majority of the sequences have a length of around 8 words. We choose to set the maximum sequence length to 10 words. This means that all sequences longer than 10 words are truncated, and all sequences shorter than 23 words are padded with a special token, '`PAD`', to make them all the same length. In addition, we discovered that some sentences are very short, therefore we choose to remove all sentences with a length less than 3 words.

Having our cleaned and tokenized text data, we build a vocabulary from the training dataset. The vocabulary is built by mapping each unique word to an integer. The vocabulary for the training dataset ended up being 10,336 words.

Then, we convert the text data to integers by mapping each word in the text data to the corresponding integer in the vocabulary and pad the sequences to the maximum sequence length. This is done for both the training, validation and test datasets. The text data is now ready to be used for a model.

4 Models

Two different models were developed for the sentiment analysis task. The first model is a Recurrent Neural Network (RNN) with Long Short-Term Memory (LSTM) cells, and the second model is a Transformer model.

4.1 Model 1: RNN-LSTM

Model Architecture We chose to begin with a Recurrent Neural Network (RNN) model using LSTM (Long Short-Term Memory) cells due to their proven effectiveness in capturing long-term dependencies in sequential data, such as text. LSTMs are particularly suited for this task because they address the vanishing gradient problem often encountered in standard RNNs.

The architecture of the model was designed to balance complexity and performance, resulting in a total of 1,117,734 trainable parameters. Key considerations behind the design are as follows:

- **Embedding Layer:** The vocabulary size of 10,336 and embedding dimension of 75 were chosen to compactly represent the input text while capturing semantic relationships between words effectively. This dimensionality ensures the embeddings are rich enough for the sentiment classification task without being computationally excessive.
- **LSTM Layer:** A single LSTM layer with 256 hidden units was selected to provide sufficient capacity to capture sequential patterns in the text. This configuration was chosen to avoid overfitting while retaining the ability to model complex dependencies in the data.
- **Dropout Layer:** A dropout rate of 50% was applied to the LSTM outputs to reduce the risk of overfitting by preventing the model from becoming overly reliant on specific neurons. This regularization technique ensures better generalization to unseen data.
- **Output Layer:** The final linear layer maps the 256 features to the 6 output classes, corresponding to the emotions in the dataset.

The specific parameters mentioned above was found using a combination of empirical testing (grid search) and theoretical considerations to achieve a balance between model complexity and performance.

Training and Evaluation The model training was conducted using the AdamW optimizer with a learning rate of 0.001 and a batch size of 16. AdamW was chosen due to its effectiveness in handling sparse gradients and its built-in weight decay mechanism, which helps prevent overfitting. The batch size of 16 strikes a balance between computational efficiency and the stability of gradient updates.

For the loss function, we used CrossEntropyLoss, which is well-suited for multi-class classification problems like this one, where the task involves predicting one of six emotion categories. To further regularize the model and reduce the risk of overfitting, an L2-norm penalty (weight decay) of 0.0001 was applied to the weights, encouraging smaller parameter values and promoting a simpler model. The model was trained for 4 epochs, as using more epochs lead to overfitting. The final test accuracy and F1-score is summarized in Table 2.




Fig. 4. Training and validation loss, accuracy, and F1-score for the RNN-LSTM model.

Label	Precision	Recall	F1-Score
sadness	0.91	0.86	0.89
joy	0.85	0.86	0.85
love	0.62	0.69	0.65
anger	0.84	0.81	0.82
fear	0.81	0.85	0.83
surprise	0.65	0.65	0.65
accuracy			0.83
weighted avg	0.84	0.83	0.83

Table 1. RNN-LSTM model performance on test set.

From the above table, we can see that the model performs well on sadness, joy, anger and fear, but not as well on love and surprise. This is likely due to the imbalanced distribution of the classes in the dataset as discussed in Section 2, see Figure 1.

4.2 Model 2: Transformer

The model is made from the solution to exercise sheet 5. Very few elements have been changed since it follows the pages in the Bishop book closely. The only big changes that has been made to the model itself is added dropout layers and positional encoding. We also experimented with randomly removing 2 tokens from each sentence to increase generalisation but it didn't work. We also tried adding 2 more linear layers to the transformer block, but this was detrimental to the accuracy so we removed them again. We chose to use a transformer model as it seems like a good architecture for working with text since it makes words weigh differently depending on which other words are present in the sentence, and with positional encoding also where they are in that sentence. This is important since the meaning of words change depending on other words, and position in sentence. We used sinusoidal position embedding since this is the positional embedding used in the Bishop book. The L in the sinusoidal position embedding is 10000. The example in the Bishop book had an L of 30, and some other examples we encountered had an L of 10000. We chose 10000 as the value for L since we encountered it multiple places.

We used cross entropy loss as our loss function since this is a good loss function for classification tasks. This is the number of heads in the multi head self attention mechanism. `d_model` should preferably be divisible by this number so each head computes the same size of input. We chose 8 heads since having too many heads can be inefficient and since we have `d_model` of 128, choosing more heads would lead to quite few dimensions per head, which would make it hard for them to capture any useful features. The `mlp_factor` controls the amount of nodes in the linear layers of the transformer block. We use an `mlp_factor` of 4. This was the default value and decreasing this value hit the performance of the model. Increasing the `mlp_factor` would increase the training time so we decided to not do that. `d_model` is the dimensions of each token. We have a vocabulary of 10336 so we made this pretty big to ensure that each word is able to be uniquely represented.

We have 6 layers of transformer blocks. According to the slides, 6 layers is a normal amount of layers and 12 layers is for very large models. We experimented with 6 and 12 layers and found only a small performance increase with 12 layers, but a significant increase in training time, so we decided that we would rather train quicker to test values for other parameters. When training the transformer models with different parameters we found out that they would quickly overfit, so we needed to generalise the model. To do this we added two dropout layers on the linear layers of the transformer blocks. The dropout chance for the first layer is 50% and for the second layer is 25%. We tested with both higher and lower dropout chances and this worked the best. We chose adam as the optimizer. From the last project our and other peoples takeaway were that adam was a good optimizer in general, and since we were not using weight decay, we don't need to use adamW. We used a learning rate of 0,0003. This was discussed as a good learning rate from the last project, and our experimenting also led us to this learning rate. We experimented a bit with using weight decay but it did not work well so we decided not to use it in our final model. When testing different parameters, we used 50 epochs. The transformer models quickly overfitted so we didn't need more epochs. We experimented with both L1 and L2 regularisation. L1 regularisation worked the best, and that is probably because it makes some weights go to 0, which indicates those words aren't important. In sentences, some words tells more than other words, so it makes sense that this would be good.




Fig. 5. Training and validation loss, accuracy, and F1-score for the transformer model.

Label	Precision	Recall	F1-Score
sadness	0.89	0.82	0.85
joy	0.81	0.88	0.85
love	0.68	0.59	0.63
anger	0.77	0.81	0.79
fear	0.80	0.80	0.80
surprise	0.67	0.56	0.61
accuracy			0.81
weighted avg	0.81	0.81	0.81

Table 2. Transformer model performance on test set.

Despite a lot of parameter tuning, we could not get the transformer model to get a better performance. We even tried adding extra linear layers in the transformer block, and removing random words from the sentences for better generalisation.

5 Analysis and Final Prediction

Of the two models, the performance of the recurrent neural network were better with an accuracy of 83% to the transformer models accuracy of 81%. Both architectures are good at predicting based on text, so its no big surprise that they are close in performance.

Since the transformer model is the worst performing, we decided to take a closer look at some failed classification cases to see why it performed as it did. These are the 7 sentences we looked at.

Number	Sentence	Transformer prediction	Actual classification
1	im updating my blog because i feel shitty	Joy	Sadness
2	i never make her separate from me because i don t ever want her to feel like i m ashamed with her	Joy	Sadness
3	i left with my bouquet of red and yellow tulips under my arm feeling slightly more optimistic than when i arrived	Sadness	Joy
4	i cant walk into a shop anywhere where i do not feel uncomfortable	Joy	Fear
5	i explain why i clung to a relationship with a boy who was in many ways immature and uncommitted despite the excitement i should have been feeling for getting accepted into the masters program at the university of virginia	Anger	Joy
6	i jest i feel grumpy tired and pre menstrual which i probably am but then again its only been a week and im about as fit as a walrus on vacation for the summer	Joy	Anger
7	i find myself in the odd position of feeling supportive of	Anger	Love

Table 3. Sentence Classifications with Transformer prediction and Actual classification

Looking at these sentences my intuition would be that the model looks at keywords to decide which category it belongs to. Some of the sentences lack any keywords that indicate what feeling they are associated with. The first sentence has a keyword 'shitty', but it might be a rare word in the training set. Sentence 1 also contains the word updating which is probably a positive word for the model, which is why it predicts joy. Sentence 6 contain the word 'vacation' which is probably why the model predicts joy for that sentence. This could indicate it is looking for keywords. If you were looking to deliberately make the model fail, we predict that misspellings, irony and figurative language would be some good techniques to do so.

To confirm our suspicions we tried to edit some of the sentences to get them correctly classified. Changing sentence 1 to "im updating my blog because i feel sad" changed the classification to sadness which is the correct one. This indicates that keywords are important. On the other hand, changing sentence 5 to "i explain why i clung to a relationship with a boy who was in many ways despite the excitement i should have been feeling for getting accepted into the masters program at the university of virginia" did not change the prediction or even the confidence which is a big surprise since the only change in that sentence is removing "immature and uncommitted" which seems like the most negative words in that sentence, so maybe the model is less keyword focused than anticipated.

6 Pretrained Model (DistilBERT)

To evaluate how our model compares to a larger pretrained model, we utilized the DistilBERT base model (uncased) from Hugging Face. DistilBERT is a distilled version of BERT, pretrained on a large corpus comprising 11,038 books and the English Wikipedia. It features a vocabulary size of 30,000 tokens and contains approximately 67 million parameters, making it significantly larger and more complex than our models. To adapt DistilBERT to our sentiment analysis task, we fine-tuned the model for 5 epochs using a small initial learning rate of 5×10^{-6} . This low learning rate ensures stable and effective fine-tuning of the pretrained weights without overfitting to our dataset. The performance of DistilBERT on our task is summarized in Table 4.

Label	Precision	Recall	F1-Score
sadness	0.92	0.93	0.92
joy	0.90	0.92	0.91
love	0.74	0.71	0.72
anger	0.90	0.89	0.89
fear	0.87	0.89	0.88
surprise	0.82	0.61	0.70
accuracy			0.89
weighted avg	0.89	0.89	0.89

Table 4. DistilBERT model performance on test set.

From the above results, we see that the pretrained model performs generally better than our two models. Notice that the performance on the less presents classes love and surprise is also lower as seen in the previous models.

7 Conclusion

We constructed two different models to complete the task, a recurrent neural network with LSTM and 256 hidden units, and a transformer model with 6 transformer layers. The recurrent neural network had the best test accuracy of 83%, while the transformer model had a test accuracy of 81%. We tried using a pretrained model distilBERT on our task to see how it would fare. It got a test accuracy of 89%. As expected our models were worse. This can be explained by the fact that distilBERT has more parameters and a much bigger training set than our models.

Overall we are satisfied with our result but it could have been improved. Initially an accuracy of 83% seems quite bad, but even the pretrained model had a test accuracy of less than 90% so compared to that, it does not seem too bad.

7.1 Individual Contributions

	Henrik Daniel Christensen	Frode Engtoft Johansen
Code	Task 1, 2, 5	Task 3, 4
Report	Section 1, 2, 3, 4.1, 6	Section 4.2, 5, 7

Table 5. Individual contributions.

A Code

Python Files

File: emotion_dataset.py

```
import torch
from torch.utils.data import Dataset

class EmotionDataset(Dataset):
    def __init__(self, data, labels):
        self.data = data
        self.labels = labels

    def __len__(self):
        return len(self.data) # return number of samples

    def __getitem__(self, idx):
        return torch.LongTensor(self.data[idx]), torch.tensor(self.labels[idx], dtype=torch.long) # return sample and l
```

File: loader.py

```
import os
import pandas as pd
from datasets import load_dataset

def loader(train_csv_path, val_csv_path, test_csv_path):
    # Check if the files exist; if not, load from the remote source
    if not (os.path.exists(train_csv_path) and os.path.exists(val_csv_path) and os.path.exists(test_csv_path)):
        print("Data files not found. Loading dataset from remote source...")

    os.makedirs("data", exist_ok=True)

    # Load the dataset from Hugging Face
    ds = load_dataset("dair-ai/emotion", "split")

    label_names = ds["train"].features["label"].names

    # Save train data
    train_data = {
        "text": ds["train"]["text"],
        "label": ds["train"]["label"],
        # Convert label indices to label names
        "label_name": [label_names[label] for label in ds["train"]["label"]]
    }
    pd.DataFrame(train_data).to_csv(train_csv_path, index=True)

    # Save validation data
    val_data = {
        "text": ds["validation"]["text"],
        "label": ds["validation"]["label"],
        "label_name": [label_names[label] for label in ds["validation"]["label"]]
    }
    pd.DataFrame(val_data).to_csv(val_csv_path, index=True)

    # Save test data
    test_data = {
        "text": ds["test"]["text"],
        "label": ds["test"]["label"],
        "label_name": [label_names[label] for label in ds["test"]["label"]]
    }
    pd.DataFrame(test_data).to_csv(test_csv_path, index=True)

    train_df = pd.read_csv(train_csv_path, index_col=0)
    val_df = pd.read_csv(val_csv_path, index_col=0)
    test_df = pd.read_csv(test_csv_path, index_col=0)

    return train_df, val_df, test_df
```

File: metrics.py

```
import os
from sklearn.metrics import (
    accuracy_score,
    f1_score,
```



```

precision_score,
recall_score,
classification_report,
confusion_matrix,
)

def compute_metrics(true_labels, predicted_labels, labels):
    accuracy = round(accuracy_score(true_labels, predicted_labels), 4)
    f1 = round(f1_score(true_labels, predicted_labels, average="weighted", zero_division=0), 4)
    precision = round(precision_score(true_labels, predicted_labels, average="weighted", zero_division=0), 4)
    recall = round(recall_score(true_labels, predicted_labels, average="weighted", zero_division=0), 4)
    class_report = classification_report(true_labels, predicted_labels, target_names=labels, zero_division=0)
    conf_matrix = confusion_matrix(true_labels, predicted_labels)
    return {
        "accuracy": accuracy,
        "f1": f1,
        "precision": precision,
        "recall": recall,
        "class_report": class_report,
        "conf_matrix": conf_matrix.tolist()
    }

def print_metrics(metrics):
    for key, value in metrics.items():
        if key == "conf_matrix":
            print("confusion matrix:")
            for row in value:
                print(row)
            print()
        elif key == "class_report":
            print("classification report:")
            print(value)
        else:
            print(f"{key}: {value}")

def save_metrics(label, metrics):
    os.makedirs("results", exist_ok=True)
    with open(f"results/{label}_metrics.txt", "w") as f:
        f.write(f"Accuracy Score: {metrics['accuracy']}\n")
        f.write(f"F1-Score: {metrics['f1']}\n")
        f.write(f"Precision: {metrics['precision']}\n")
        f.write(f"Recall: {metrics['recall']}\n\n")
        f.write("Classification Report:\n")
        f.write(metrics['class_report'] + "\n")
        f.write("Confusion Matrix:\n")
        for row in metrics['conf_matrix']:
            f.write(str(row) + "\n")

```

File: plot_scores.py

```

import matplotlib.pyplot as plt
from matplotlib.ticker import MaxNLocator

def plot_scores(results, label):
    _, axes = plt.subplots(1, 3, figsize=(9, 3)) # Adjust the layout for 3 subplots

    epochs = range(1, results["num_epochs"] + 1)

    # Plot loss
    axes[0].plot(epochs, results["train_losses"], marker="o", markersize=3, label="Train Loss")
    axes[0].plot(epochs, results["val_losses"], marker="o", markersize=3, linestyle="--", label="Val Loss")
    axes[0].set_title(f"{label}: Loss")
    axes[0].set_xlabel("Epoch")
    axes[0].set_ylabel("Loss")
    axes[0].legend(loc="best")

    # Convert accuracy and F1-score to percentages
    train_accuracies = [acc * 100 for acc in results["train_accuracies"]]
    val_accuracies = [acc * 100 for acc in results["val_accuracies"]]
    train_f1_scores = [f1 * 100 for f1 in results["train_f1_scores"]]
    val_f1_scores = [f1 * 100 for f1 in results["val_f1_scores"]]

    # Plot accuracy
    axes[1].plot(epochs, train_accuracies, marker="o", markersize=3, label="Train Accuracy")
    axes[1].plot(epochs, val_accuracies, marker="o", markersize=3, linestyle="--", label="Val Accuracy")
    axes[1].set_title(f"{label}: Accuracy")

```

```

axes[1].set_xlabel("Epoch")
axes[1].set_ylabel("Accuracy (%)")
axes[1].legend(loc="best")

# Plot F1-score
axes[2].plot(epochs, train_f1_scores, marker="o", markersize=3, label="Train F1-Score")
axes[2].plot(epochs, val_f1_scores, marker="o", markersize=3, linestyle="--", label="Val F1-Score")
axes[2].set_title(f"{label}: F1-Score")
axes[2].set_xlabel("Epoch")
axes[2].set_ylabel("F1-Score (%)")
axes[2].legend(loc="best")

# Set the number of ticks on the x and y axes for all plots
for ax in axes:
    ax.yaxis.set_major_locator(MaxNLocator(nbins=4))
    ax.xaxis.set_major_locator(MaxNLocator(nbins=4))

# Adjust layout and display the plots
plt.tight_layout()
plt.show()

```

File: predict.py

```

import os
import torch
import pandas as pd
from metrics import compute_metrics, print_metrics, save_metrics

def decode_tokens(token_ids, reverse_vocab):
    words = [reverse_vocab[token] for token in token_ids if token in reverse_vocab]
    return " ".join(words)

def predict(label, model, device, loader, label_map, reverse_vocab):
    model.to(device)
    model.eval()

    # Store predictions
    predictions = []
    true_labels = []
    predicted_labels = []

    # Ensure the results directory exists
    os.makedirs("results", exist_ok=True)

    with torch.no_grad():
        for batch in loader:
            sentences, labels = batch
            labels = labels.to(device)

            # Forward pass
            outputs = model(sentences.to(device))
            probabilities = torch.softmax(outputs, dim=1)
            predicted = torch.argmax(probabilities, dim=1)

            # Collect true and predicted labels for metrics
            true_labels.extend(labels.cpu().numpy())
            predicted_labels.extend(predicted.cpu().numpy())

            # Convert tokenized tensors back to text using decode_tokens()
            decoded_sentences = [
                decode_tokens(sentence.tolist(), reverse_vocab)
                for sentence in sentences
            ]

        for i, decoded_sentence in enumerate(decoded_sentences):
            true_label = label_map[labels[i].item()]
            pred_label = label_map[predicted[i].item()]
            confidence = probabilities[i][predicted[i].item()].item()
            correct = true_label == pred_label

            predictions.append({
                "Sentence": decoded_sentence,
                "Correct": correct,
                "True Label": true_label,
                "Predicted Label": pred_label,
                "Confidence (%)": f"{confidence * 100:.2f}"
            })

```

```

    })

    # Save predictions to a CSV file
    predictions_df = pd.DataFrame(predictions)
    predictions_file = f"results/{label}_predictions.csv"
    predictions_df.to_csv(predictions_file, index=False)

    # Compute metrics
    metrics = compute_metrics(true_labels, predicted_labels, label_map.values())
    print_metrics(metrics)
    save_metrics(label, metrics)

```

File: train_model.py

```

import os
import time
import numpy as np
from collections import Counter

from torch import nn
import torch
import torch.optim as optim
from torch.nn.utils import clip_grad_norm_
from torchinfo import summary

from sklearn.metrics import accuracy_score, f1_score

from metrics import compute_metrics, print_metrics, save_metrics

def log_undefined(predicted_labels, labels):
    counts = Counter(predicted_labels)
    for idx, label in enumerate(labels):
        if counts[idx] == 0:
            print(f"Warning: No predictions for label '{label}' (index {idx}).")

def log_undefined(predicted_labels, labels):
    """Logs warnings for labels that are not predicted."""
    counts = Counter(predicted_labels)
    for idx, label in enumerate(labels):
        if counts[idx] == 0:
            print(f"Warning: No predictions for label '{label}' (index {idx}).")

def train_model(
    label,
    model,
    train_loader,
    val_loader,
    label_map,
    device,
    optimizer_type="Adam",
    learning_rate=0.001,
    momentum=0.9,
    weight_decay=0.0,
    step_size=None,
    gamma=0.5,
    reg_type=None,
    reg_lambda=0.0,
    num_epochs=30,
    grad_clip=0.0,
):
    """Trains a PyTorch model and logs metrics for each epoch."""
    # Move the model to the device
    model = model.to(device)

    # Define the loss function
    criterion = nn.CrossEntropyLoss()

    # Select optimizer
    if optimizer_type == "SGD":
        optimizer = optim.SGD(
            model.parameters(), lr=learning_rate, momentum=momentum, weight_decay=weight_decay
        )
    elif optimizer_type == "Adam":
        optimizer = optim.Adam(

```

```

        model.parameters(), lr=learning_rate, weight_decay=weight_decay
    )
    elif optimizer_type == "AdamW":
        optimizer = optim.AdamW(
            model.parameters(), lr=learning_rate, weight_decay=weight_decay
        )
    else:
        raise ValueError(f"Unknown optimizer type: {optimizer_type}")

    # Learning rate scheduler
    scheduler = None
    if step_size is not None and gamma is not None:
        scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=step_size, gamma=gamma)

    # Metrics storage
    train_losses, train_accuracies, train_f1_scores = [], [], []
    val_losses, val_accuracies, val_f1_scores = [], [], []

    total_start_time = time.time()

    for epoch in range(num_epochs):
        start_time = time.time()

        # Training phase
        model.train()

        epoch_total_train_loss = 0.0
        epoch_total_train_samples = 0
        epoch_train_true_labels = []
        epoch_train_predicted_labels = []

        for inputs, targets in train_loader:
            inputs, targets = inputs.to(device), targets.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            # Apply regularization
            if reg_lambda > 0.0 and reg_type is not None:
                if reg_type == "L1":
                    l1_norm = sum(param.abs().sum() for param in model.parameters())
                    loss += reg_lambda * l1_norm
                elif reg_type == "L2":
                    l2_norm = sum(param.pow(2).sum() for param in model.parameters())
                    loss += reg_lambda * l2_norm

            loss.backward()

            if grad_clip > 0:
                clip_grad_norm_(model.parameters(), grad_clip)

            optimizer.step()

            epoch_total_train_loss += loss.item() * inputs.size(0)
            epoch_total_train_samples += inputs.size(0)

            # Collect true labels and predictions
            _, predicted = torch.max(outputs, dim=1)
            epoch_train_true_labels.extend(targets.cpu().numpy())
            epoch_train_predicted_labels.extend(predicted.cpu().numpy())

        # Calculate training metrics
        avg_epoch_train_loss = round(epoch_total_train_loss / epoch_total_train_samples, 4)
        epoch_train_accuracy = round(accuracy_score(epoch_train_true_labels, epoch_train_predicted_labels), 4)
        epoch_train_f1 = round(f1_score(epoch_train_true_labels, epoch_train_predicted_labels, average='weighted'), 4)
        train_losses.append(avg_epoch_train_loss)
        train_accuracies.append(epoch_train_accuracy)
        train_f1_scores.append(epoch_train_f1)

        # Validation phase
        model.eval()

        epoch_total_val_loss = 0.0
        epoch_total_val_samples = 0
        all_val_true_labels = []
        all_val_predicted_labels = []

```

```

with torch.no_grad():
    for inputs, targets in val_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = model(inputs)

        avg_epoch_val_loss = criterion(outputs, targets)
        epoch_total_val_loss += avg_epoch_val_loss.item() * inputs.size(0)
        epoch_total_val_samples += inputs.size(0)

        # Collect true labels and predictions
        _, predicted = torch.max(outputs, dim=1)
        all_val_true_labels.extend(targets.cpu().numpy())
        all_val_predicted_labels.extend(predicted.cpu().numpy())

    # Calculate validation metrics
    avg_epoch_val_loss = round(epoch_total_val_loss / epoch_total_val_samples, 4)
    epoch_val_accuracy = round(accuracy_score(all_val_true_labels, all_val_predicted_labels), 4)
    epoch_val_f1 = round(f1_score(all_val_true_labels, all_val_predicted_labels, average='weighted'), 4)
    val_losses.append(avg_epoch_val_loss)
    val_accuracies.append(epoch_val_accuracy)
    val_f1_scores.append(epoch_val_f1)

    # Update learning rate
    if scheduler:
        scheduler.step()

    epoch_duration = round(time.time() - start_time)
    print(
        f"Epoch {epoch + 1}/{num_epochs} ({epoch_duration}s) | "
        f"Train: loss {avg_epoch_train_loss}, acc {epoch_train_accuracy*100:.2f}%, f1 {epoch_train_f1*100:.2f}% | "
        f"Val: loss {avg_epoch_val_loss}, acc {epoch_val_accuracy*100:.2f}%, f1 {epoch_val_f1*100:.2f}%"
    )

    # Log undefined predictions
    log_undefined(all_val_predicted_labels, label_map.values())

    # Total training time
    total_training_time = round(time.time() - total_start_time)
    print(f"Total Training Time: {total_training_time}s\n")

    # Save model summary and metrics
    os.makedirs("models", exist_ok=True)
    with open(f"models/{label}.txt", "w", encoding="utf-8") as f:
        f.write(str(summary(model, verbose=0)))

    # Compute and save metrics
    metrics = compute_metrics(all_val_true_labels, all_val_predicted_labels, label_map.values())
    print_metrics(metrics)
    save_metrics(label, metrics)

    # Save model state
    torch.save(model.state_dict(), f"models/{label}.pth")

    # Return training history
    return {
        "num_epochs": num_epochs,
        "train_losses": train_losses,
        "train_accuracies": train_accuracies,
        "train_f1_scores": train_f1_scores,
        "val_losses": val_losses,
        "val_accuracies": val_accuracies,
        "val_f1_scores": val_f1_scores,
    }

```

File: transformer_model.py

```

import torch
import torch.nn as nn
import torch.nn.functional as F
import math
import random

# the self attention is just like described in deep learning by Bishop, so i will not change it.
class SelfAttention(nn.Module):
    def __init__(self, d_model, d_key):
        super().__init__()
        # Three separate linear layers for the queries, keys, and values

```

```

        self.w_q = nn.Linear(d_model, d_key)
        self.w_k = nn.Linear(d_model, d_key)
        self.w_v = nn.Linear(d_model, d_model)

    def forward(self, x):
        q = self.w_q(x)
        k = self.w_k(x)
        v = self.w_v(x)
        # Compute the attention weights
        a = q @ k.transpose(-2, -1) / (k.shape[-1] ** 0.5)
        a = F.softmax(a, dim=-1)
        # Apply the attention weights
        z = a @ v
        return z
# Same as in book, shouldn't need any change
class MultiHeadSelfAttention(nn.Module):
    def __init__(self, d_model, d_key, n_heads):
        super().__init__()
        self.heads = nn.ModuleList([SelfAttention(d_model, d_key) for _ in range(n_heads)])
        # Down projection back to model dimension
        self.w_o = nn.Linear(n_heads * d_model, d_model)

    def forward(self, x):
        return self.w_o(torch.cat([h(x) for h in self.heads], dim=-1))
# maybe change siLU activation function?
class TransformerBlock(nn.Module):
    def __init__(self, d_model, d_key, n_heads, dropout1, dropout2, mlp_factor=4, ):
        super().__init__()
        # We need to init two layer norms because they have parameters
        self.ln1 = nn.LayerNorm(d_model)
        self.attn = MultiHeadSelfAttention(d_model, d_key, n_heads)
        self.ln2 = nn.LayerNorm(d_model)
        # a feedforward module
        if dropout1 > 0:
            self.mlp = nn.Sequential(
                nn.Linear(d_model, mlp_factor * d_model),
                nn.Dropout(p = dropout1),
                nn.SiLU(), # Swish activation function, f(x) = x * sigmoid(x)
                nn.Linear(mlp_factor * d_model, d_model),
                nn.Dropout(p = dropout2)
            )
        else:
            self.mlp = nn.Sequential(
                nn.Linear(d_model, mlp_factor * d_model),
                nn.SiLU(), # Swish activation function, f(x) = x * sigmoid(x)
                nn.Linear(mlp_factor * d_model, d_model),
                nn.SiLU(),
                nn.Linear(d_model, mlp_factor * d_model),
                nn.SiLU(),
                nn.Linear(d_model, mlp_factor * d_model),
            )

    def forward(self, x):
        # Residual connections and pre-layernorm
        x = x + self.attn(self.ln1(x))
        x = x + self.mlp(self.ln2(x))
        return x

class TransformerClassifier(nn.Module):
    def __init__(self, n_embeds, n_classes, d_model=256, d_key=64, n_heads=2, mlp_factor=4, n_layers=2, device = "cpu",
                 ):
        super().__init__()
        self.device = device
        self.d_model = d_model
        self.token_embedding = nn.Embedding(n_embeds, d_model)
        self.transformer_model = nn.Sequential(*[TransformerBlock(d_model, d_key, n_heads, dropout1, dropout2, mlp_factor)
                                                for _ in range(n_layers)])
        self.final_layer_norm = nn.LayerNorm(d_model)
        self.classifier = nn.Sequential(nn.Linear(d_model, d_model), nn.SiLU(), nn.Linear(d_model, n_classes))

    def sinusoidalPositionEncoding(self, input):
        # create empty matrix
        r_n_matrix = torch.empty((input.size(dim=1), self.d_model))
        r_n_matrix = r_n_matrix.to(self.device)
        # fill all areas of empty matrix
        for n in range(input.size(dim=1)):
            for i in range(self.d_model):
                if i % 2 == 0:
                    r_n_matrix[n, i] = math.sin(n / 10000 ** (i / self.d_model))
                if i % 2 == 1:
                    r_n_matrix[n, i] = math.cos(n / 10000 ** (i / self.d_model))
        # add with input
        input_hat = input + r_n_matrix
        # return modified input

```

```
    return input_hat

def forward(self, x):
    e = self.token_embedding(x)
    # sinusoidal positional encoding
    s = self.sinusoidalPositionEncoding(e)
    h = self.transformer_model(s)
    h = h.mean(dim=1) # Average pooling on the sequence dimension
    y = self.classifier(self.final_layer_norm(h))
    return y
```

B Notebook

notebook

November 29, 2024

0.1 Libraries

```
[4]: %load_ext autoreload
      %autoreload 2

[5]: from collections import Counter
      from collections import defaultdict
      import itertools
      import json
      import matplotlib.pyplot as plt
      import pandas as pd
      import numpy as np
      import nltk
      from nltk.tokenize import RegexpTokenizer
      from nltk.corpus import stopwords
      from nltk.stem import PorterStemmer
      nltk.download('wordnet') # downloads WordNet data
      nltk.download('omw-1.4') # downloads additional wordnet data for lemmatization
      nltk.download('stopwords') # downloads stopwords (if not already downloaded)
      from sklearn.metrics import (
          accuracy_score,
          f1_score,
      )
      from wordcloud import WordCloud
      import torch
      import torch.nn as nn
      from torch.utils.data import DataLoader
      from datasets import Dataset
      from transformers import (
          AutoTokenizer,
          AutoModelForSequenceClassification,
          Trainer,
          TrainingArguments,
      )

      from emotion_dataset import EmotionDataset
      from loader import loader
      from train_model import train_model
      from plot_scores import plot_scores
      from predict import predict
      from predict_on_fly import predict_on_fly
      from metrics import compute_metrics, print_metrics, save_metrics
      import transformer_model

[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\difj6\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to
[nltk_data] C:\Users\difj6\AppData\Roaming\nltk_data...
[nltk_data] Package omw-1.4 is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\difj6\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
c:\Users\difj6\OneDrive - Syddansk Universitet\Documents\Uni\7. semester\DM873
Deep learning\Project2\.venv\lib\site-packages\tqdm\auto.py:21: TqdmWarning:
IProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```


0.2 Device

```
[6]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
      print(f"Using {device}")
```

Using cuda

0.3 Load and save dataset

```
[7]: train_csv_path = "data/train.csv"
      val_csv_path = "data/val.csv"
      test_csv_path = "data/test.csv"

      train_df, val_df, test_df = loader(train_csv_path, val_csv_path, test_csv_path)
```

1 Task 1: Data Preparation

1.1 Data set

```
[8]: print(f"# train sentences: {len(train_df)}")
      print(f"# validation sentences: {len(val_df)}")
      print(f"# test sentences: {len(test_df)}")

      print(train_df)
```

```
# train sentences: 16000
# validation sentences: 2000
# test sentences: 2000
```

	text	label	label_name
0	i didnt feel humiliated	0	sadness
1	i can go from feeling so hopeless to so damned...	0	sadness
2	im grabbing a minute to post i feel greedy wrong	3	anger
3	i am ever feeling nostalgic about the fireplac...	2	love
4	i am feeling grouchy	3	anger
...
15995	i just had a very brief time in the beanbag an...	0	sadness
15996	i am now turning and i feel pathetic that i am...	0	sadness
15997	i feel strong and good overall	1	joy
15998	i feel like this was such a rude comment and i...	3	anger
15999	i know a lot but i feel so stupid because i ca...	0	sadness

[16000 rows x 3 columns]

1.2 Step 1: Dataset Preparation

1.2.1 Label distribution

```
[9]: # Calculate label distributions and percentages for training and validation sets
      def calculate_distribution(df):
          label_distribution = df["label_name"].value_counts().reset_index()
          label_distribution.columns = ["label_name", "count"]
          label_distribution["percentage"] = round((label_distribution["count"] / label_distribution["count"].
          ↳sum()) * 100, 2)
          return label_distribution

      # Calculate distributions
      train_distribution = calculate_distribution(train_df)
      val_distribution = calculate_distribution(val_df)

      # Merge with label mapping for alignment
      label_map_df = train_df[["label", "label_name"]].drop_duplicates().sort_values("label")
      label_map = dict(zip(label_map_df["label"], label_map_df["label_name"]))
      train_labels_df = label_map_df.merge(train_distribution, on="label_name", how="left")
      val_labels_df = label_map_df.merge(val_distribution, on="label_name", how="left")

      # Print training label mapping and distribution
      num_classes = len(label_map_df)
      print(f"Number of classes: {num_classes}")
```

```

print("Training Label Mapping and Distribution:")
print(train_labels_df.to_string(index=False))

# Plot side-by-side
fig, axes = plt.subplots(1, 2, figsize=(6, 3), sharey=True)

# Training set
axes[0].bar(train_labels_df["label_name"], train_labels_df["percentage"], color="skyblue",
            edgecolor="black")
axes[0].set_title("Training Label Distribution")
axes[0].set_xlabel("Labels")
axes[0].set_ylabel("Percentage (%)")
axes[0].tick_params(axis='x', rotation=45)

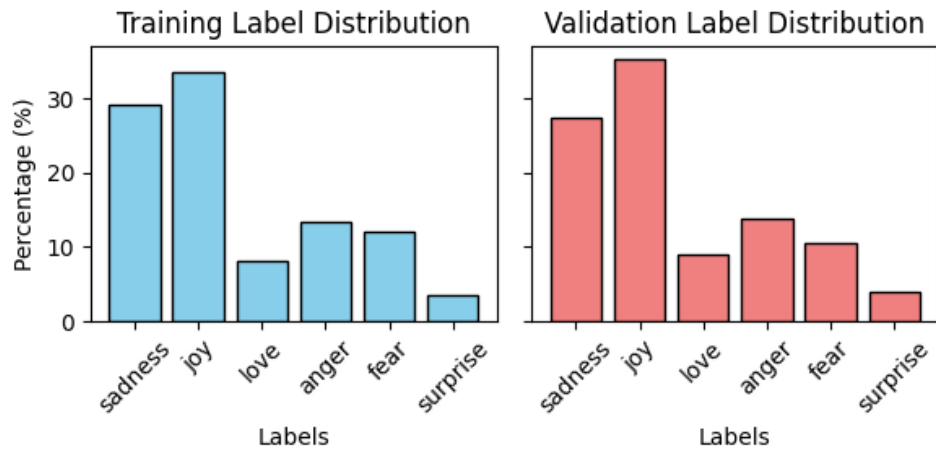
# Validation set
axes[1].bar(val_labels_df["label_name"], val_labels_df["percentage"], color="lightcoral", edgecolor="black")
axes[1].set_title("Validation Label Distribution")
axes[1].set_xlabel("Labels")
axes[1].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

```

Number of classes: 6
 Training Label Mapping and Distribution:

label	label_name	count	percentage
0	sadness	4666	29.16
1	joy	5362	33.51
2	love	1304	8.15
3	anger	2159	13.49
4	fear	1937	12.11
5	surprise	572	3.58



1.3 Step 2: Tokenizing

1.3.1 Tokenizer

```

[10]: tokenizer = RegexpTokenizer(r"[a-zA-Z]+|([?'\"]+)" # sequence that don't match the pattern act as
      ↪ separators.
example_sentence = "This?.is,a:cu123stom;tokenization example!<"
example_tokens = tokenizer.tokenize(example_sentence)
print(example_tokens)

```

```
['This', '?', 'is', 'a', 'cu', 'stom', 'tokenization', 'example', '!']
```

1.3.2 Tokenize each split

```
[11]: train_df["tokens"] = train_df["text"].str.lower().apply(tokenizer.tokenize)
      val_df["tokens"] = val_df["text"].str.lower().apply(tokenizer.tokenize)
      test_df["tokens"] = test_df["text"].str.lower().apply(tokenizer.tokenize)

      train_vocab = set(token for tokens in train_df["tokens"] for token in tokens)

      print(f"# words in train vocab: {len(train_vocab)}")
```

```
# words in train vocab: 15212
```

1.3.3 Word frequency

```
[12]: def get_top_words_per_class(tokens_in, top_n=10):
      tokens_by_class = defaultdict(list)
      for tokens, label in zip(tokens_in, train_df["label_name"]):
          tokens_by_class[label].append(tokens)
      tokens_by_class = dict(tokens_by_class)
      results = []
      for label_name, tokens in tokens_by_class.items():
          flat_tokens = list(itertools.chain.from_iterable(tokens))
          most_common = Counter(flat_tokens).most_common(top_n)
          for word, count in most_common:
              results.append({"Label_name": label_name, "Word": word, "Count": count})
      return pd.DataFrame(results)

      print(get_top_words_per_class(train_df["tokens"], top_n=10).to_string(index=False))
```

Label_name	Word	Count
sadness	i	7635
sadness	feel	3299
sadness	and	2692
sadness	to	2335
sadness	the	2155
sadness	a	1656
sadness	feeling	1523
sadness	of	1422
sadness	that	1299
sadness	my	1245
anger	i	3576
anger	feel	1459
anger	and	1258
anger	to	1162
anger	the	1109
anger	a	791
anger	feeling	721
anger	that	705
anger	of	630
anger	my	573
love	i	2120
love	feel	929
love	and	902
love	to	860
love	the	780
love	a	571
love	of	482
love	that	460
love	my	399
love	feeling	378
surprise	i	927
surprise	feel	356
surprise	and	354
surprise	the	335
surprise	to	267
surprise	a	256

surprise	that	212
surprise	feeling	209
surprise	of	191
surprise	my	163
fear	i	3083
fear	feel	1212
fear	to	1116
fear	and	1110
fear	the	1000
fear	a	806
fear	feeling	742
fear	of	614
fear	that	531
fear	my	525
joy	i	8518
joy	feel	3928
joy	and	3273
joy	to	3232
joy	the	2991
joy	a	2120
joy	that	1905
joy	of	1651
joy	feeling	1539
joy	my	1378

1.3.4 Data cleaning (remove stopwords and stem)

```
[13]: stemmer = PorterStemmer()
stop_words = set(stopwords.words("english"))

def preprocess_tokens(tokens):
    return [stemmer.stem(word) for word in tokens if word not in stop_words]

train_df["processed_tokens"] = train_df["tokens"].apply(preprocess_tokens)
val_df["processed_tokens"] = val_df["tokens"].apply(preprocess_tokens)
test_df["processed_tokens"] = test_df["tokens"].apply(preprocess_tokens)

print(get_top_words_per_class(train_df["processed_tokens"], top_n=10).to_string(index=False))
```

Label_name	Word	Count
sadness	feel	4994
sadness	like	881
sadness	im	683
sadness	know	297
sadness	get	284
sadness	realli	276
sadness	time	271
sadness	make	245
sadness	want	244
sadness	go	235
anger	feel	2261
anger	like	391
anger	im	342
anger	get	175
anger	time	140
anger	want	133
anger	irrit	128
anger	realli	124
anger	know	122
anger	hate	113
love	feel	1406
love	like	366
love	love	277
love	im	193
love	support	103
love	realli	92
love	know	89
love	want	89

love	time	82
love	care	82
surprise	feel	601
surprise	amaz	107
surprise	like	92
surprise	im	91
surprise	impress	63
surprise	overwhelm	58
surprise	weird	57
surprise	surpris	56
surprise	curiou	54
surprise	funni	49
fear	feel	2025
fear	im	322
fear	like	264
fear	littl	149
fear	go	139
fear	know	136
fear	bit	118
fear	want	113
fear	time	110
fear	get	107
joy	feel	5674
joy	like	1023
joy	im	799
joy	make	381
joy	time	334
joy	get	322
joy	go	315
joy	realli	309
joy	want	272
joy	know	261

1.3.5 Frequency of words

```
[14]: word_freq = Counter(itertools.chain.from_iterable(train_df["processed_tokens"]))
      pd.DataFrame(word_freq.items(), columns=["Word", "Count"]).sort_values(by="Count", ascending=False).
      ↪to_csv("results/word_frequencies.csv", index=False)
      print(pd.DataFrame(word_freq.most_common(100), columns=["Word", "Count"]).to_string(index=False))
```

Word	Count
feel	16961
like	3017
im	2430
get	981
time	979
realli	942
know	938
make	935
go	882
want	867
love	805
littl	736
think	736
day	675
thing	672
peopl	664
one	647
would	646
even	600
still	598
ive	587
life	555
way	528
need	521
bit	521
someth	514
much	496

dont	482
work	471
could	453
say	450
start	445
look	423
see	419
back	414
tri	410
good	408
pretti	392
right	357
alway	356
come	351
help	342
friend	340
also	337
year	336
today	332
use	326
take	317
around	315
person	303
cant	301
made	296
hate	285
well	279
though	274
happi	274
didnt	272
got	271
write	270
live	268
felt	266
lot	264
never	264
thought	263
hope	261
someon	259
find	259
everi	254
quit	250
read	246
less	246
sure	240
enough	238
week	236
give	234
mani	232
kind	230
home	227
away	226
support	224
long	222
ever	221
anyth	220
actual	220
talk	215
better	213
keep	212
left	211
let	210
everyth	210
without	209
rememb	209
last	207
care	205
tell	205

```

world    205
wonder   204
sometim  201
new       199
http      199

```

1.3.6 Remove additional words

```

[15]: additional_words_to_remove = ["feel", "realli", "im", "know", "also", "http"]

def remove_additional_words(tokens):
    return [word for word in tokens if word not in additional_words_to_remove]

train_df["processed_tokens"] = train_df["processed_tokens"].apply(remove_additional_words)
val_df["processed_tokens"] = val_df["processed_tokens"].apply(remove_additional_words)
test_df["processed_tokens"] = test_df["processed_tokens"].apply(remove_additional_words)

train_df["processed_text"] = train_df["processed_tokens"].apply(" ".join)
val_df["processed_text"] = val_df["processed_tokens"].apply(" ".join)
test_df["processed_text"] = test_df["processed_tokens"].apply(" ".join)

print(get_top_words_per_class(train_df["processed_tokens"], top_n=10).to_string(index=False))



```

Label_name	Word	Count
sadness	like	881
sadness	get	284
sadness	time	271
sadness	make	245
sadness	want	244
sadness	go	235
sadness	day	224
sadness	thing	221
sadness	ive	217
sadness	think	212
anger	like	391
anger	get	175
anger	time	140
anger	want	133
anger	irrit	128
anger	hate	113
anger	thing	109
anger	make	108
anger	go	108
anger	think	105
love	like	366
love	love	277
love	support	103
love	want	89
love	time	82
love	care	82
love	long	72
love	one	70
love	get	70
love	sweet	69
surprise	amaz	107
surprise	like	92
surprise	impress	63
surprise	overwhelm	58
surprise	weird	57
surprise	surpris	56
surprise	curiou	54
surprise	funni	49
surprise	strang	46
surprise	shock	46
fear	like	264
fear	littl	149
fear	go	139
fear	bit	118


```
plt.xlabel("Number of Words")
plt.yticks([])
plt.show()

# Plot distribution of lengths
plt.figure(figsize=(2.3, 1.5))
plt.hist(train_lengths, bins=30, color="skyblue", edgecolor="black")
plt.title("Training Text Lengths")
plt.xlabel("Number of Words")
plt.ylabel("Count")
plt.show()
```

Length range for train: from 1 to 34 words
Mean length for train: 8 words
Standard deviation for train: 5



1.3.9 Set max and min length

```
[18]: # Set max length for the model. If sentence is longer, truncate it. If shorter, pad it.
# Set min length to remove very short sentences from the training set.
max_length = 10
min_length = 3
```

```
print(f"# train sentences before filtering: {len(train_df)}")
train_df = train_df[train_df["processed_tokens"].apply(len) >= min_length]
print(f"# train sentences after filtering: {len(train_df)}")
```

train sentences before filtering: 16000
train sentences after filtering: 14331

1.4 Step 3: Build a vocabulary

```
[19]: vocab = {"<PAD>": 0, "<UNK>": 1}
      for tokens in train_df["processed_tokens"]:
          for token in tokens:
              if token not in vocab:
                  vocab[token] = len(vocab)

      vocab_size = len(vocab)
      print(f"Vocabulary size: {vocab_size}")

      reverse_vocab = {v: k for k, v in vocab.items()}
```

Vocabulary size: 10336

1.5 Step 4: Encode all texts with the vocabulary

```
[20]: def encode(tokens): # i.e. words to integers
      return [vocab[token] if token in vocab else 1 for token in tokens]

      train_df["encoded"] = train_df["processed_tokens"].apply(encode)
      val_df["encoded"] = val_df["processed_tokens"].apply(encode)
      test_df["encoded"] = test_df["processed_tokens"].apply(encode)
```

1.6 Step 5: Maximum sequence length

```
[21]: def pad(sequence):
      return sequence[:max_length] + [0] * (max_length - len(sequence))

      train_df["padded"] = train_df["encoded"].apply(pad)
      val_df["padded"] = val_df["encoded"].apply(pad)
      test_df["padded"] = test_df["encoded"].apply(pad)
```

2 Task 2: RNN model

2.1 Model class

```
[72]: class RNN_model(nn.Module):
      def __init__(self, type, vocab_size, embedding_dim, hidden_size, num_classes, padding_idx=0,
      ↪ num_layers=1, dropout_rnn=0, dropout_fc=0):
          super(RNN_model, self).__init__()
          # embedding_dim: size of each embedding vector
          # hidden_size: number of features in the hidden state
          # num_layers: number of recurrent layers
          # bias: introduces a bias
          # batch_first: input and output tensors are provided as (batch, seq, feature)
          # dropout: if non-zero, introduces a dropout layer on the outputs of each RNN layer except the last
          ↪ layer

          self.embedding = nn.Embedding(num_embeddings=vocab_size, embedding_dim=embedding_dim,
          ↪ padding_idx=padding_idx)
          if type == "RNN":
              self.rnn = nn.RNN(input_size=embedding_dim, hidden_size=hidden_size, num_layers=num_layers,
          ↪ bias=True, batch_first=True, dropout=dropout_rnn, nonlinearity="tanh")
          elif type == "GRU":
              self.rnn = nn.GRU(input_size=embedding_dim, hidden_size=hidden_size, num_layers=num_layers,
          ↪ bias=True, batch_first=True, dropout=dropout_rnn)
          elif type == "LSTM":
              self.rnn = nn.LSTM(input_size=embedding_dim, hidden_size=hidden_size, num_layers=num_layers,
          ↪ bias=True, batch_first=True, dropout=dropout_rnn)
          self.fc = nn.Linear(in_features=hidden_size, out_features=num_classes)
          self.dropout = nn.Dropout(p=dropout_fc)

      def forward(self, x):
          x = self.embedding(x)
          x, _ = self.rnn(x)
```

```

x = x[:, -1, :] # extract last hidden state for each sequence
x = self.dropout(x) # apply dropout to the last hidden state
x = self.fc(x) # pass last hidden state through the fully connected layer
return x

```

2.2 Hyper parameter tuning

```

[73]: train_dataset = EmotionDataset(train_df["padded"].tolist(), train_df["label"].tolist())
      val_dataset = EmotionDataset(val_df["padded"].tolist(), val_df["label"].tolist())
      test_dataset = EmotionDataset(test_df["padded"].tolist(), test_df["label"].tolist())

[ ]: def grid_search(vocab_size, num_classes, train_dataset, val_dataset, label_map, device):
      param_grid = {
          # The best hyperparameters found is commented out
          "type": ["LSTM", "GRU", "RNN"], # "LSTM"
          "embedding_dim": [100, 75], # 75
          "hidden_size": [512, 256], # 256
          "layers": [1, 2], # 1
          "dropout_rnn": [0.0, 0.2], # 0.0
          "dropout_fc": [0.4, 0.6], # 0.4
          "learning_rate": [0.001, 0.0005], # 0.0001
          "reg_lambda": [0.0001, 0.00005], # 0.0001
          "batch_size": [16, 32], # 16
      }

      # Generate all combinations of hyperparameters
      keys, values = zip(*param_grid.items())
      configs = [dict(zip(keys, v)) for v in itertools.product(*values)]
      total_configs = len(configs)

      best_val_f1_score = 0
      results_list = []
      for i, config in enumerate(configs):
          label = f"model_{i}"
          print(f"Training model: {label} ({i+1}/{total_configs})")
          print(json.dumps(config, indent=4))

          train_loader = DataLoader(train_dataset, batch_size=config["batch_size"], shuffle=True)
          val_loader = DataLoader(val_dataset, batch_size=config["batch_size"], shuffle=False)

          # Model
          model = RNN_model(
              type=config["type"],
              vocab_size=vocab_size,
              embedding_dim=config["embedding_dim"],
              hidden_size=config["hidden_size"],
              num_classes=num_classes,
              num_layers=config["layers"],
              dropout_rnn=config["dropout_rnn"],
              dropout_fc=config["dropout_fc"],
          )

          # Train
          results = train_model(
              label=label,
              model=model,
              train_loader=train_loader,
              val_loader=val_loader,
              label_map=label_map,
              device=device,
              optimizer_type="AdamW",
              learning_rate=config["learning_rate"],
              reg_type="L2",
              reg_lambda=config["reg_lambda"],
              num_epochs=10
          )

```

```

# Track the best model
current_val_f1_score = max(results["val_f1_scores"])
if current_val_f1_score > best_val_f1_score:
    best_val_f1_score = current_val_f1_score
    best_model_label = label
    print(f"New best model found: {best_model_label} with val f1 score: {best_val_f1_score:.4f}")

# Add results to list
results["config"] = config
results_list.append(results)

# Save results
with open("results/grid_search_results.json", "w") as f:
    json.dump(results_list, f, indent=4)

print(f"Best model found: {best_model_label}")

grid_search(vocab_size, num_classes, train_dataset, val_dataset, label_map, device)

```

2.3 Best RNN model

```

[75]: label = "best_rnn_model"
      best_batch_size = 16

      train_loader = DataLoader(train_dataset, batch_size=best_batch_size, shuffle=True)
      val_loader = DataLoader(val_dataset, batch_size=best_batch_size, shuffle=False)

      best_rnn_model = RNN_model(
          type="LSTM",
          vocab_size=vocab_size,
          embedding_dim=75,
          hidden_size=256,
          num_classes=num_classes,
          num_layers=1,
          dropout_rnn=0.0,
          dropout_fc=0.4,
      )

      best_rnn_results = train_model(
          label=label,
          model=best_rnn_model,
          train_loader=train_loader,
          val_loader=val_loader,
          label_map=label_map,
          device=device,
          optimizer_type="AdamW",
          learning_rate=0.001,
          reg_type="L2",
          reg_lambda=0.0001,
          num_epochs=4
      )

      plot_scores(best_rnn_results, label)

```

```

Epoch 1/4 (10s) | Train: loss 43.7021, acc 40.13%, f1 34.61% | Val: loss 1.1907,
acc 57.05%, f1 54.13%
Epoch 2/4 (7s) | Train: loss 13.1735, acc 69.10%, f1 68.43% | Val: loss 0.6395,
acc 80.35%, f1 80.08%
Epoch 3/4 (7s) | Train: loss 5.0673, acc 83.09%, f1 82.97% | Val: loss 0.5178,
acc 83.10%, f1 82.92%
Epoch 4/4 (7s) | Train: loss 2.5631, acc 86.90%, f1 86.83% | Val: loss 0.551,
acc 82.20%, f1 82.18%
Total Training Time: 32s

```

```

accuracy: 0.822
f1: 0.8218
precision: 0.8257

```

```

recall: 0.822
classification report:


```

	precision	recall	f1-score	support
sadness	0.80	0.91	0.85	550
joy	0.89	0.83	0.86	704
love	0.69	0.69	0.69	178
anger	0.86	0.76	0.80	275
fear	0.79	0.80	0.79	212
surprise	0.75	0.77	0.76	81
accuracy			0.82	2000
macro avg	0.79	0.79	0.79	2000
weighted avg	0.83	0.82	0.82	2000

```

confusion matrix:
[502, 15, 7, 8, 16, 2]
[54, 581, 42, 10, 10, 7]
[15, 29, 122, 6, 3, 3]
[41, 14, 3, 208, 8, 1]
[16, 6, 4, 9, 169, 8]
[3, 5, 0, 2, 9, 62]

```



3 Task 3: Transformer model

```

[76]: batch_size = 32
train_dataset = EmotionDataset(train_df["padded"].tolist(), train_df["label"].tolist())
val_dataset = EmotionDataset(val_df["padded"].tolist(), val_df["label"].tolist())
test_dataset = EmotionDataset(test_df["padded"].tolist(), test_df["label"].tolist())
train_labels = train_df["label"].tolist()
val_labels = val_df["label"].tolist()
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
n_embeds = 10336
num_epochs = 50

```

```

[81]: config = {
    "label": "transformer1",
    "d_key": 64,
    "n_heads": 8,
    "mlp_factor": 4,
    "d_model": 128,
    "n_layers": 6,
    "dropout1": 0.5,
    "dropout2": 0.25,
    "optimizer_type": "Adam",
    "learning_rate": 0.0003,
    "weight_decay": 0,

```

```

        "reg_type": "L1",
        "reg_lambda": 1e-4,
    }
    print(f"Training model: {config['label']}")
    model = transformer_model.TransformerClassifier(n_embeds=n_embeds, n_classes=6, d_model=config["d_model"],
    ↪ d_key=config["d_key"], n_heads=config["n_heads"], mlp_factor=config["mlp_factor"],
    ↪ n_layers=config["n_layers"], device = device, dropout1=config["dropout1"], dropout2=config["dropout2"])
    results = train_model(
        label=config["label"],
        model=model,
        train_loader=train_loader,
        val_loader=val_loader,
        label_map=label_map,
        device=device,
        optimizer_type=config["optimizer_type"],
        learning_rate=config["learning_rate"],
        weight_decay=config["weight_decay"],
        reg_type=config["reg_type"],
        reg_lambda=config["reg_lambda"],
        num_epochs=num_epochs,
    )
    plot_scores(results, config["label"])

```

```

Training model: transformer1
Epoch 1/50 (89s) | Train: loss 102.0006, acc 33.86%, f1 23.71% | Val: loss
1.5793, acc 32.10%, f1 23.49%
Epoch 2/50 (76s) | Train: loss 82.9869, acc 36.89%, f1 27.59% | Val: loss
1.5281, acc 40.85%, f1 31.31%
Epoch 3/50 (76s) | Train: loss 69.0236, acc 41.38%, f1 33.03% | Val: loss
1.4376, acc 45.90%, f1 38.03%
Epoch 4/50 (80s) | Train: loss 57.1575, acc 46.12%, f1 41.04% | Val: loss
1.3466, acc 50.80%, f1 46.88%
Epoch 5/50 (77s) | Train: loss 47.2545, acc 52.08%, f1 48.79% | Val: loss
1.1978, acc 54.85%, f1 50.44%
Epoch 6/50 (80s) | Train: loss 39.011, acc 59.49%, f1 57.57% | Val: loss 1.062,
acc 62.00%, f1 60.50%
Epoch 7/50 (79s) | Train: loss 32.3509, acc 65.00%, f1 63.95% | Val: loss
0.9193, acc 68.25%, f1 67.81%
Epoch 8/50 (77s) | Train: loss 26.9274, acc 70.02%, f1 69.34% | Val: loss
0.8478, acc 70.40%, f1 69.82%
Epoch 9/50 (76s) | Train: loss 22.5769, acc 73.58%, f1 73.13% | Val: loss
0.8064, acc 71.35%, f1 71.28%
Epoch 10/50 (76s) | Train: loss 19.0561, acc 76.75%, f1 76.42% | Val: loss
0.712, acc 74.90%, f1 74.65%
Epoch 11/50 (76s) | Train: loss 16.2046, acc 78.24%, f1 77.94% | Val: loss
0.6758, acc 77.20%, f1 77.05%
Epoch 12/50 (76s) | Train: loss 13.8662, acc 79.91%, f1 79.69% | Val: loss
0.6526, acc 77.35%, f1 76.76%
Epoch 13/50 (76s) | Train: loss 11.9736, acc 81.45%, f1 81.25% | Val: loss
0.6205, acc 79.95%, f1 79.66%
Epoch 14/50 (77s) | Train: loss 10.4343, acc 82.03%, f1 81.87% | Val: loss
0.6295, acc 78.65%, f1 78.57%
Epoch 15/50 (77s) | Train: loss 9.1926, acc 82.99%, f1 82.84% | Val: loss
0.5959, acc 79.95%, f1 79.67%
Epoch 16/50 (76s) | Train: loss 8.1612, acc 83.69%, f1 83.55% | Val: loss
0.5867, acc 80.30%, f1 80.02%
Epoch 17/50 (77s) | Train: loss 7.3331, acc 83.83%, f1 83.73% | Val: loss
0.5698, acc 81.00%, f1 80.82%
Epoch 18/50 (76s) | Train: loss 6.6518, acc 84.61%, f1 84.49% | Val: loss
0.6136, acc 79.55%, f1 79.36%
Epoch 19/50 (76s) | Train: loss 6.104, acc 84.82%, f1 84.73% | Val: loss 0.5767,
acc 80.55%, f1 80.34%
Epoch 20/50 (76s) | Train: loss 5.6224, acc 85.53%, f1 85.45% | Val: loss
0.5924, acc 80.15%, f1 79.88%
Epoch 21/50 (76s) | Train: loss 5.2473, acc 85.56%, f1 85.48% | Val: loss
0.6482, acc 78.85%, f1 78.85%
Epoch 22/50 (76s) | Train: loss 4.9044, acc 86.11%, f1 86.02% | Val: loss
0.5677, acc 80.80%, f1 80.52%

```

```

Epoch 23/50 (76s) | Train: loss 4.6268, acc 86.14%, f1 86.07% | Val: loss
0.5594, acc 82.25%, f1 82.19%
Epoch 24/50 (76s) | Train: loss 4.3677, acc 86.46%, f1 86.37% | Val: loss
0.5706, acc 81.20%, f1 80.95%
Epoch 25/50 (78s) | Train: loss 4.1486, acc 86.97%, f1 86.90% | Val: loss
0.5459, acc 81.90%, f1 81.75%
Epoch 26/50 (76s) | Train: loss 3.9806, acc 86.69%, f1 86.62% | Val: loss 0.542,
acc 82.40%, f1 82.38%
Epoch 27/50 (77s) | Train: loss 3.7892, acc 87.52%, f1 87.46% | Val: loss 0.542,
acc 82.30%, f1 81.97%
Epoch 28/50 (76s) | Train: loss 3.6444, acc 87.35%, f1 87.29% | Val: loss
0.5416, acc 82.55%, f1 82.39%
Epoch 29/50 (76s) | Train: loss 3.4955, acc 87.80%, f1 87.74% | Val: loss
0.5703, acc 81.15%, f1 80.88%
Epoch 30/50 (77s) | Train: loss 3.389, acc 87.71%, f1 87.65% | Val: loss 0.546,
acc 82.00%, f1 81.85%
Epoch 31/50 (76s) | Train: loss 3.2559, acc 88.16%, f1 88.10% | Val: loss
0.5646, acc 82.30%, f1 81.96%
Epoch 32/50 (76s) | Train: loss 3.1771, acc 87.64%, f1 87.59% | Val: loss
0.5844, acc 81.15%, f1 81.21%
Epoch 33/50 (76s) | Train: loss 3.0501, acc 88.59%, f1 88.54% | Val: loss
0.5427, acc 82.45%, f1 82.40%
Epoch 34/50 (76s) | Train: loss 2.9593, acc 88.60%, f1 88.55% | Val: loss
0.5736, acc 81.90%, f1 81.82%
Epoch 35/50 (76s) | Train: loss 2.8801, acc 88.61%, f1 88.56% | Val: loss
0.5598, acc 82.25%, f1 82.16%
Epoch 36/50 (77s) | Train: loss 2.8071, acc 88.83%, f1 88.78% | Val: loss
0.5595, acc 82.15%, f1 82.04%
Epoch 37/50 (77s) | Train: loss 2.7384, acc 88.84%, f1 88.79% | Val: loss
0.6097, acc 81.05%, f1 80.46%
Epoch 38/50 (76s) | Train: loss 2.659, acc 89.01%, f1 88.96% | Val: loss 0.5831,
acc 81.10%, f1 80.82%
Epoch 39/50 (77s) | Train: loss 2.5879, acc 88.94%, f1 88.89% | Val: loss
0.5631, acc 82.45%, f1 82.42%
Epoch 40/50 (77s) | Train: loss 2.5379, acc 89.44%, f1 89.40% | Val: loss 0.589,
acc 82.45%, f1 82.54%
Epoch 41/50 (77s) | Train: loss 2.4701, acc 89.35%, f1 89.31% | Val: loss
0.5726, acc 82.80%, f1 82.71%
Epoch 42/50 (76s) | Train: loss 2.4347, acc 89.32%, f1 89.28% | Val: loss
0.5533, acc 83.55%, f1 83.54%
Epoch 43/50 (77s) | Train: loss 2.3524, acc 89.97%, f1 89.93% | Val: loss
0.5814, acc 82.50%, f1 82.40%
Epoch 44/50 (76s) | Train: loss 2.3024, acc 90.38%, f1 90.34% | Val: loss
0.5854, acc 83.15%, f1 83.07%
Epoch 45/50 (77s) | Train: loss 2.2664, acc 89.99%, f1 89.95% | Val: loss
0.5644, acc 83.30%, f1 83.14%
Epoch 46/50 (76s) | Train: loss 2.2204, acc 90.38%, f1 90.34% | Val: loss
0.5628, acc 82.60%, f1 82.43%
Epoch 47/50 (76s) | Train: loss 2.167, acc 90.64%, f1 90.61% | Val: loss 0.606,
acc 81.65%, f1 81.72%
Epoch 48/50 (76s) | Train: loss 2.1392, acc 90.64%, f1 90.62% | Val: loss
0.5831, acc 83.00%, f1 82.94%
Epoch 49/50 (76s) | Train: loss 2.0786, acc 91.10%, f1 91.07% | Val: loss
0.5902, acc 82.90%, f1 82.65%
Epoch 50/50 (76s) | Train: loss 2.056, acc 90.95%, f1 90.92% | Val: loss 0.594,
acc 82.80%, f1 82.62%
Total Training Time: 3843s

```

accuracy: 0.828

f1: 0.8262

precision: 0.8267

recall: 0.828


classification report:

	precision	recall	f1-score	support
sadness	0.88	0.85	0.87	550
joy	0.83	0.89	0.86	704
love	0.74	0.60	0.66	178

anger	0.81	0.82	0.81	275
fear	0.79	0.81	0.80	212
surprise	0.74	0.68	0.71	81
accuracy			0.83	2000
macro avg	0.80	0.78	0.79	2000
weighted avg	0.83	0.83	0.83	2000

confusion matrix:

```
[470, 34, 7, 18, 18, 3]
[23, 628, 24, 17, 5, 7]
[8, 51, 107, 8, 3, 1]
[16, 17, 4, 225, 12, 1]
[14, 12, 2, 6, 171, 7]
[2, 12, 0, 4, 8, 55]
```



Train multiple:

```
[ ]: configurations = [
    #heads 16
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 16,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L2",
        "reg_lambda": 1e-4,
    },
    # mlp_factor 2
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 2,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L2",
    },
]
```



```

        "reg_lambda": 1e-4,
    },
    # 0,0003 learning rate
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.0003,
        "weight_decay": 0,
        "reg_type": "L2",
        "reg_lambda": 1e-4,
    },
    # L1 reg
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L1",
        "reg_lambda": 1e-4,
    },
    # More dropout
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.6,
        "dropout2": 0.3,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L2",
        "reg_lambda": 1e-4,
    },
]

for config in configurations:
    print(f"Training model: {config['label']}")
    model = transformer_model.TransformerClassifier(n_embeddings=n_embeddings, n_classes=6,
    ↪ d_model=config["d_model"], d_key=config["d_key"], n_heads=config["n_heads"],
    ↪ mlp_factor=config["mlp_factor"], n_layers=config["n_layers"], device = device,
    ↪ dropout1=config["dropout1"], dropout2=config["dropout2"])
    results = train_model(
        label=config["label"],
        model=model,
        train_loader=train_loader,
        val_loader=val_loader,
        label_map=label_map,
        device=device,
        optimizer_type=config["optimizer_type"],
        learning_rate=config["learning_rate"],
        weight_decay=config["weight_decay"],

```

```

    reg_type=config["reg_type"],
    reg_lambda=config["reg_lambda"],
    num_epochs=num_epochs,
)
plot_scores(results, config["label"])

```

4 Task 4: Analysis

```

[ ]: test_loader = DataLoader(test_dataset, batch_size=best_batch_size, shuffle=True)

predict("best_rnn_model", best_rnn_model, device, test_loader, label_map, reverse_vocab)

```

```

accuracy: 0.8245
f1: 0.8248
precision: 0.8261
recall: 0.8245
classification report:

```

	precision	recall	f1-score	support
sadness	0.85	0.90	0.87	581
joy	0.87	0.82	0.84	695
love	0.64	0.70	0.67	159
anger	0.83	0.81	0.82	275
fear	0.81	0.81	0.81	224
surprise	0.64	0.62	0.63	66
accuracy			0.82	2000
macro avg	0.77	0.78	0.78	2000
weighted avg	0.83	0.82	0.82	2000

```

confusion matrix:
[520, 31, 4, 12, 12, 2]
[35, 571, 52, 20, 8, 9]
[16, 23, 112, 4, 2, 2]
[20, 20, 4, 223, 7, 1]
[17, 6, 2, 8, 182, 9]
[4, 7, 0, 1, 13, 41]

```

```

[82]: test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True)

predict("best_rnn_model", model, device, test_loader, label_map, reverse_vocab)

```

```

accuracy: 0.813
f1: 0.8115
precision: 0.8126
recall: 0.813
classification report:

```

	precision	recall	f1-score	support
sadness	0.89	0.82	0.85	581
joy	0.81	0.88	0.85	695
love	0.68	0.59	0.63	159
anger	0.77	0.81	0.79	275
fear	0.80	0.80	0.80	224
surprise	0.67	0.56	0.61	66
accuracy			0.81	2000
macro avg	0.77	0.75	0.76	2000
weighted avg	0.81	0.81	0.81	2000

```

confusion matrix:
[479, 52, 6, 28, 13, 3]
[21, 614, 29, 15, 9, 7]
[7, 46, 94, 11, 0, 1]
[18, 19, 7, 222, 9, 0]
[15, 10, 2, 10, 180, 7]

```

```
[0, 13, 0, 3, 13, 37]
```

```
[ ]: config = {
    "label": "transformer1",
    "d_key": 64,
    "n_heads": 8,
    "mlp_factor": 4,
    "d_model": 128,
    "n_layers": 6,
    "dropout1": 0.5,
    "dropout2": 0.25,
    "optimizer_type": "Adam",
    "learning_rate": 0.0003,
    "weight_decay": 0,
    "reg_type": "L1",
    "reg_lambda": 1e-4,
}
model = transformer_model.TransformerClassifier(n_embeds=n_embeds, n_classes=6, d_model=config["d_model"],
↪ d_key=config["d_key"], n_heads=config["n_heads"], mlp_factor=config["mlp_factor"],
↪ n_layers=config["n_layers"], device = device, dropout1=config["dropout1"], dropout2=config["dropout2"])
model.load_state_dict(torch.load('models/transformer1.pth'))
predict_on_fly(model, tokenizer, vocab, device, label_map, max_length)
```

5 Task 5: Pre-trained model (transfer learning)

```
[ ]: model_name = "distilbert-base-uncased"
num_epochs = 5
learning_rate = 5e-6
max_length = 32 # different max length as we are using raw data to be tokenized using the pretrained
↪ tokenizer
batch_size = 32
label2id = {v: k for k, v in label_map.items()}

# Metric function
def pretrained_evaluation(predictions):
    preds = predictions.predictions.argmax(-1) # get predicted labels
    labels = predictions.label_ids
    accuracy = accuracy_score(labels, preds)
    f1 = f1_score(labels, preds, average="weighted")
    return {"accuracy": accuracy, "f1": f1}

# Load tokenizer
tokenizer_pretrained = AutoTokenizer.from_pretrained(model_name)

# Preprocessing function
def tokenize(batch):
    return tokenizer_pretrained(batch["text"], padding="max_length", truncation=True, max_length=max_length)

# Convert pandas DataFrame to Hugging Face Dataset
train_dataset = Dataset.from_pandas(train_df)
val_dataset = Dataset.from_pandas(val_df)
test_dataset = Dataset.from_pandas(test_df)

# Tokenize datasets
train_dataset = train_dataset.map(tokenize, batched=True)
val_dataset = val_dataset.map(tokenize, batched=True)
test_dataset = test_dataset.map(tokenize, batched=True)

# Set format for PyTorch
train_dataset = train_dataset.rename_column("label", "labels")
val_dataset = val_dataset.rename_column("label", "labels")
test_dataset = test_dataset.rename_column("label", "labels")
train_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
val_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
test_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
```

```

# Load pre-trained model
model = AutoModelForSequenceClassification.from_pretrained(
    model_name,
    num_labels=num_classes,
    id2label=label_map,
    label2id=label2id
)

# Training arguments
training_args = TrainingArguments(
    output_dir="results",
    eval_strategy="epoch", # evaluate at the end of each epoch
    save_strategy="epoch", # save checkpoints after every epoch
    logging_strategy="epoch",
    learning_rate=learning_rate,
    per_device_train_batch_size=batch_size,
    per_device_eval_batch_size=64,
    weight_decay=0.01,
    num_train_epochs=num_epochs,
    load_best_model_at_end=True, # load best model at the end of training
    metric_for_best_model="f1", # specify metric to monitor
    save_total_limit=1, # keep only the best checkpoint
)

# Define Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    compute_metrics=pretrained_evaluation, # compute metrics during evaluation
)

# Train model
trainer.train()

# Evaluate best model
best_results = trainer.evaluate()
print(f"Best Model Evaluation Results:\n {best_results}")

# Save best fine-tuned model
fine_tuned_model_name = "distilbert_finetuned"
trainer.save_model(f"models/{fine_tuned_model_name}")
tokenizer_pretrained.save_pretrained(f"models/{fine_tuned_model_name}")

```

```

[32]: # Load fine-tuned model
tokenizer_new = AutoTokenizer.from_pretrained(f"models/{fine_tuned_model_name}")
model_new = AutoModelForSequenceClassification.from_pretrained(f"models/{fine_tuned_model_name}",
    ↪ num_labels=num_classes)
trainer = Trainer(model=model_new)

# Predict on test set
predictions = trainer.predict(test_dataset)

# Extract logits and compute predicted labels
logits = torch.tensor(predictions.predictions) # convert logits to a PyTorch tensor
predicted_labels = torch.argmax(logits, dim=1).numpy() # convert to numpy array for sklearn metrics

# Evaluate predictions
metrics = compute_metrics(test_df["label"], predicted_labels, label_map.values())
print_metrics(metrics)
save_metrics(fine_tuned_model_name, metrics)

```

```
100%| 250/250 [00:01<00:00, 152.01it/s]
```

```
accuracy: 0.8885
```

```
f1: 0.8871
```

```

precision: 0.8869
recall: 0.8885
classification report:
      precision    recall  f1-score   support

   sadness      0.92      0.93      0.93      581
     joy      0.89      0.92      0.91      695
     love      0.74      0.70      0.72      159
     anger      0.90      0.90      0.90      275
     fear      0.90      0.87      0.88      224
   surprise      0.80      0.61      0.69       66

 accuracy              0.89      2000
  macro avg      0.86      0.82      0.84      2000
 weighted avg      0.89      0.89      0.89      2000

confusion matrix:
[543, 20, 3, 11, 4, 0]
[14, 641, 32, 6, 1, 1]
[6, 39, 111, 3, 0, 0]
[14, 7, 2, 247, 5, 0]
[11, 2, 0, 7, 195, 9]
[4, 9, 1, 0, 12, 40]

```