

notebook

November 29, 2024

0.1 Libraries

```
[4]: %load_ext autoreload
      %autoreload 2
```

```
[5]: from collections import Counter
      from collections import defaultdict
      import itertools
      import json
      import matplotlib.pyplot as plt
      import pandas as pd
      import numpy as np
      import nltk
      from nltk.tokenize import RegexpTokenizer
      from nltk.corpus import stopwords
      from nltk.stem import PorterStemmer
      nltk.download('wordnet') # downloads WordNet data
      nltk.download('omw-1.4') # downloads additional wordnet data for lemmatization
      nltk.download('stopwords') # downloads stopwords (if not already downloaded)
      from sklearn.metrics import (
          accuracy_score,
          f1_score,
      )
      from wordcloud import WordCloud
      import torch
      import torch.nn as nn
      from torch.utils.data import DataLoader
      from datasets import Dataset
      from transformers import (
          AutoTokenizer,
          AutoModelForSequenceClassification,
          Trainer,
          TrainingArguments,
      )

      from emotion_dataset import EmotionDataset
      from loader import loader
      from train_model import train_model
      from plot_scores import plot_scores
      from predict import predict
      from predict_on_fly import predict_on_fly
      from metrics import compute_metrics, print_metrics, save_metrics
      import transformer_model
```

```
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\difj6\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to
[nltk_data] C:\Users\difj6\AppData\Roaming\nltk_data...
[nltk_data] Package omw-1.4 is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\difj6\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
c:\Users\difj6\OneDrive - Syddansk Universitet\Documents\Uni\7. semester\DM873
Deep learning\Project2\venv\lib\site-packages\tqdm\auto.py:21: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

0.2 Device

```
[6]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
     print(f"Using {device}")
```

Using cuda

0.3 Load and save dataset

```
[7]: train_csv_path = "data/train.csv"
     val_csv_path = "data/val.csv"
     test_csv_path = "data/test.csv"

     train_df, val_df, test_df = loader(train_csv_path, val_csv_path, test_csv_path)
```

1 Task 1: Data Preparation

1.1 Data set

```
[8]: print(f"# train sentences: {len(train_df)}")
     print(f"# validation sentences: {len(val_df)}")
     print(f"# test sentences: {len(test_df)}")

     print(train_df)
```

```
# train sentences: 16000
# validation sentences: 2000
# test sentences: 2000
```

	text	label	label_name
0	i didnt feel humiliated	0	sadness
1	i can go from feeling so hopeless to so damned...	0	sadness
2	im grabbing a minute to post i feel greedy wrong	3	anger
3	i am ever feeling nostalgic about the fireplac...	2	love
4	i am feeling grouchy	3	anger
...
15995	i just had a very brief time in the beanbag an...	0	sadness
15996	i am now turning and i feel pathetic that i am...	0	sadness
15997	i feel strong and good overall	1	joy
15998	i feel like this was such a rude comment and i...	3	anger
15999	i know a lot but i feel so stupid because i ca...	0	sadness

[16000 rows x 3 columns]

1.2 Step 1: Dataset Preparation

1.2.1 Label distribution

```
[9]: # Calculate label distributions and percentages for training and validation sets
def calculate_distribution(df):
    label_distribution = df["label_name"].value_counts().reset_index()
    label_distribution.columns = ["label_name", "count"]
    label_distribution["percentage"] = round((label_distribution["count"] / label_distribution["count"].
    ↪sum()) * 100, 2)
    return label_distribution

# Calculate distributions
train_distribution = calculate_distribution(train_df)
val_distribution = calculate_distribution(val_df)

# Merge with label mapping for alignment
label_map_df = train_df[["label", "label_name"]].drop_duplicates().sort_values("label")
label_map = dict(zip(label_map_df["label"], label_map_df["label_name"]))
train_labels_df = label_map_df.merge(train_distribution, on="label_name", how="left")
val_labels_df = label_map_df.merge(val_distribution, on="label_name", how="left")

# Print training label mapping and distribution
num_classes = len(label_map_df)
print(f"Number of classes: {num_classes}")
```

```

print("Training Label Mapping and Distribution:")
print(train_labels_df.to_string(index=False))

# Plot side-by-side
fig, axes = plt.subplots(1, 2, figsize=(6, 3), sharey=True)

# Training set
axes[0].bar(train_labels_df["label_name"], train_labels_df["percentage"], color="skyblue",
            edgecolor="black")
axes[0].set_title("Training Label Distribution")
axes[0].set_xlabel("Labels")
axes[0].set_ylabel("Percentage (%)")
axes[0].tick_params(axis='x', rotation=45)

# Validation set
axes[1].bar(val_labels_df["label_name"], val_labels_df["percentage"], color="lightcoral", edgecolor="black")
axes[1].set_title("Validation Label Distribution")
axes[1].set_xlabel("Labels")
axes[1].tick_params(axis='x', rotation=45)


plt.tight_layout()
plt.show()

```

Number of classes: 6

Training Label Mapping and Distribution:

label	label_name	count	percentage
0	sadness	4666	29.16
1	joy	5362	33.51
2	love	1304	8.15
3	anger	2159	13.49
4	fear	1937	12.11
5	surprise	572	3.58



1.3 Step 2: Tokenizing

1.3.1 Tokenizer

```

[10]: tokenizer = RegexpTokenizer(r"[a-zA-Z]+|([!?'\" ]+)" # sequence that don't match the pattern act as
      ↪ separators.
example_sentence = "This?.is,a:cu123stom;tokenization example!<"
example_tokens = tokenizer.tokenize(example_sentence)
print(example_tokens)

```

```
['This', '?', 'is', 'a', 'cu', 'stom', 'tokenization', 'example', '!']
```

1.3.2 Tokenize each split

```
[11]: train_df["tokens"] = train_df["text"].str.lower().apply(tokenizer.tokenize)
      val_df["tokens"] = val_df["text"].str.lower().apply(tokenizer.tokenize)
      test_df["tokens"] = test_df["text"].str.lower().apply(tokenizer.tokenize)

      train_vocab = set(token for tokens in train_df["tokens"] for token in tokens)

      print(f"# words in train vocab: {len(train_vocab)}")
```

```
# words in train vocab: 15212
```

1.3.3 Word frequency

```
[12]: def get_top_words_per_class(tokens_in, top_n=10):
      tokens_by_class = defaultdict(list)
      for tokens, label in zip(tokens_in, train_df["label_name"]):
          tokens_by_class[label].append(tokens)
      tokens_by_class = dict(tokens_by_class)
      results = []
      for label_name, tokens in tokens_by_class.items():
          flat_tokens = list(itertools.chain.from_iterable(tokens))
          most_common = Counter(flat_tokens).most_common(top_n)
          for word, count in most_common:
              results.append({"Label_name": label_name, "Word": word, "Count": count})
      return pd.DataFrame(results)

      print(get_top_words_per_class(train_df["tokens"], top_n=10).to_string(index=False))
```

Label_name	Word	Count
sadness	i	7635
sadness	feel	3299
sadness	and	2692
sadness	to	2335
sadness	the	2155
sadness	a	1656
sadness	feeling	1523
sadness	of	1422
sadness	that	1299
sadness	my	1245
anger	i	3576
anger	feel	1459
anger	and	1258
anger	to	1162
anger	the	1109
anger	a	791
anger	feeling	721
anger	that	705
anger	of	630
anger	my	573
love	i	2120
love	feel	929
love	and	902
love	to	860
love	the	780
love	a	571
love	of	482
love	that	460
love	my	399
love	feeling	378
surprise	i	927
surprise	feel	356
surprise	and	354
surprise	the	335
surprise	to	267
surprise	a	256

surprise	that	212
surprise	feeling	209
surprise	of	191
surprise	my	163
fear	i	3083
fear	feel	1212
fear	to	1116
fear	and	1110
fear	the	1000
fear	a	806
fear	feeling	742
fear	of	614
fear	that	531
fear	my	525
joy	i	8518
joy	feel	3928
joy	and	3273
joy	to	3232
joy	the	2991
joy	a	2120
joy	that	1905
joy	of	1651
joy	feeling	1539
joy	my	1378

1.3.4 Data cleaning (remove stopwords and stem)

```
[13]: stemmer = PorterStemmer()
stop_words = set(stopwords.words("english"))

def preprocess_tokens(tokens):
    return [stemmer.stem(word) for word in tokens if word not in stop_words]

train_df["processed_tokens"] = train_df["tokens"].apply(preprocess_tokens)
val_df["processed_tokens"] = val_df["tokens"].apply(preprocess_tokens)
test_df["processed_tokens"] = test_df["tokens"].apply(preprocess_tokens)

print(get_top_words_per_class(train_df["processed_tokens"], top_n=10).to_string(index=False))
```

Label_name	Word	Count
sadness	feel	4994
sadness	like	881
sadness	im	683
sadness	know	297
sadness	get	284
sadness	realli	276
sadness	time	271
sadness	make	245
sadness	want	244
sadness	go	235
anger	feel	2261
anger	like	391
anger	im	342
anger	get	175
anger	time	140
anger	want	133
anger	irrit	128
anger	realli	124
anger	know	122
anger	hate	113
love	feel	1406
love	like	366
love	love	277
love	im	193
love	support	103
love	realli	92
love	know	89
love	want	89

love	time	82
love	care	82
surprise	feel	601
surprise	amaz	107
surprise	like	92
surprise	im	91
surprise	impress	63
surprise	overwhelm	58
surprise	weird	57
surprise	surpris	56
surprise	curiou	54
surprise	funni	49
fear	feel	2025
fear	im	322
fear	like	264
fear	littl	149
fear	go	139
fear	know	136
fear	bit	118
fear	want	113
fear	time	110
fear	get	107
joy	feel	5674
joy	like	1023
joy	im	799
joy	make	381
joy	time	334
joy	get	322
joy	go	315
joy	realli	309
joy	want	272
joy	know	261

1.3.5 Frequency of words

```
[14]: word_freq = Counter(itertools.chain.from_iterable(train_df["processed_tokens"]))
      pd.DataFrame(word_freq.items(), columns=["Word", "Count"]).sort_values(by="Count", ascending=False).
      ↪to_csv("results/word_frequencies.csv", index=False)
      print(pd.DataFrame(word_freq.most_common(100), columns=["Word", "Count"]).to_string(index=False))
```

Word	Count
feel	16961
like	3017
im	2430
get	981
time	979
realli	942
know	938
make	935
go	882
want	867
love	805
littl	736
think	736
day	675
thing	672
peopl	664
one	647
would	646
even	600
still	598
ive	587
life	555
way	528
need	521
bit	521
someth	514
much	496

dont	482
work	471
could	453
say	450
start	445
look	423
see	419
back	414
tri	410
good	408
pretti	392
right	357
always	356
come	351
help	342
friend	340
also	337
year	336
today	332
use	326
take	317
around	315
person	303
cant	301
made	296
hate	285
well	279
though	274
happi	274
didnt	272
got	271
write	270
live	268
felt	266
lot	264
never	264
thought	263
hope	261
someon	259
find	259
everi	254
quit	250
read	246
less	246
sure	240
enough	238
week	236
give	234
mani	232
kind	230
home	227
away	226
support	224
long	222
ever	221
anyth	220
actual	220
talk	215
better	213
keep	212
left	211
let	210
everyth	210
without	209
rememb	209
last	207
care	205
tell	205

world	205
wonder	204
sometim	201
new	199
http	199

1.3.6 Remove additional words

```
[15]: additional_words_to_remove = ["feel", "realli", "im", "know", "also", "http"]

def remove_additional_words(tokens):
    return [word for word in tokens if word not in additional_words_to_remove]

train_df["processed_tokens"] = train_df["processed_tokens"].apply(remove_additional_words)
val_df["processed_tokens"] = val_df["processed_tokens"].apply(remove_additional_words)
test_df["processed_tokens"] = test_df["processed_tokens"].apply(remove_additional_words)

train_df["processed_text"] = train_df["processed_tokens"].apply(" ".join)
val_df["processed_text"] = val_df["processed_tokens"].apply(" ".join)
test_df["processed_text"] = test_df["processed_tokens"].apply(" ".join)

print(get_top_words_per_class(train_df["processed_tokens"], top_n=10).to_string(index=False))
```

Label_name	Word	Count
sadness	like	881
sadness	get	284
sadness	time	271
sadness	make	245
sadness	want	244
sadness	go	235
sadness	day	224
sadness	thing	221
sadness	ive	217
sadness	think	212
anger	like	391
anger	get	175
anger	time	140
anger	want	133
anger	irrit	128
anger	hate	113
anger	thing	109
anger	make	108
anger	go	108
anger	think	105
love	like	366
love	love	277
love	support	103
love	want	89
love	time	82
love	care	82
love	long	72
love	one	70
love	get	70
love	sweet	69
surprise	amaz	107
surprise	like	92
surprise	impress	63
surprise	overwhelm	58
surprise	weird	57
surprise	surpris	56
surprise	curiou	54
surprise	funni	49
surprise	strang	46
surprise	shock	46
fear	like	264
fear	littl	149
fear	go	139
fear	bit	118

fear	want	113
fear	time	110
fear	get	107
fear	make	105
fear	think	94
fear	peopl	90
joy	like	1023
joy	make	381
joy	time	334
joy	get	322
joy	go	315
joy	want	272
joy	love	257
joy	day	241
joy	think	233
joy	one	211

1.3.7 WordCloud

```
[16]: wordcloud = WordCloud(width=400, height=200, background_color="white").generate(" ".join(list(itertools.
    ↪ chain.from_iterable(train_df["processed_tokens"]))))
plt.figure(figsize=(5, 3))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off")
plt.show()
```



1.3.8 Sentence length distribution

```
[17]: train_lengths = [len(tokens) for tokens in train_df["processed_tokens"]]
      mean_length = np.mean(train_lengths)
      std_dev = np.std(train_lengths)



      print(f"Length range for train: from {min(train_lengths)} to {max(train_lengths)} words")
      print(f"Mean length for train: {mean_length:.0f} words")
      print(f"Standard deviation for train: {std_dev:.0f}")

      # Plot boxplot
      plt.figure(figsize=(3, 1.2))
      plt.boxplot(
          train_lengths,
          vert=False,
          patch_artist=True,
          boxprops=dict(facecolor="skyblue", linewidth=2), # larger, colored box
          whiskerprops=dict(linewidth=2), # Thicker whiskers
          medianprops=dict(color="red", linewidth=2), # highlight the median
      )
      plt.title("Training Text Lengths")
```

```
plt.xlabel("Number of Words")
plt.yticks([])
plt.show()

# Plot distribution of lengths
plt.figure(figsize=(2.3, 1.5))
plt.hist(train_lengths, bins=30, color="skyblue", edgecolor="black")
plt.title("Training Text Lengths")
plt.xlabel("Number of Words")
plt.ylabel("Count")
plt.show()
```

Length range for train: from 1 to 34 words
Mean length for train: 8 words
Standard deviation for train: 5



1.3.9 Set max and min length

```
[18]: # Set max length for the model. If sentence is longer, truncate it. If shorter, pad it.
# Set min length to remove very short sentences from the training set.
max_length = 10
min_length = 3

print(f"# train sentences before filtering: {len(train_df)}")
train_df = train_df[train_df["processed_tokens"].apply(len) >= min_length]
print(f"# train sentences after filtering: {len(train_df)}")

# train sentences before filtering: 16000
# train sentences after filtering: 14331
```

1.4 Step 3: Build a vocabulary

```
[19]: vocab = {"<PAD>": 0, "<UNK>": 1}
      for tokens in train_df["processed_tokens"]:
          for token in tokens:
              if token not in vocab:
                  vocab[token] = len(vocab)

      vocab_size = len(vocab)
      print(f"Vocabulary size: {vocab_size}")

      reverse_vocab = {v: k for k, v in vocab.items()}
```

Vocabulary size: 10336

1.5 Step 4: Encode all texts with the vocabulary

```
[20]: def encode(tokens): # i.e. words to integers
      return [vocab[token] if token in vocab else 1 for token in tokens]

      train_df["encoded"] = train_df["processed_tokens"].apply(encode)
      val_df["encoded"] = val_df["processed_tokens"].apply(encode)
      test_df["encoded"] = test_df["processed_tokens"].apply(encode)
```

1.6 Step 5: Maximum sequence length

```
[21]: def pad(sequence):
      return sequence[:max_length] + [0] * (max_length - len(sequence))

      train_df["padded"] = train_df["encoded"].apply(pad)
      val_df["padded"] = val_df["encoded"].apply(pad)
      test_df["padded"] = test_df["encoded"].apply(pad)
```

2 Task 2: RNN model

2.1 Model class

```
[72]: class RNN_model(nn.Module):
      def __init__(self, type, vocab_size, embedding_dim, hidden_size, num_classes, padding_idx=0,
      ↪ num_layers=1, dropout_rnn=0, dropout_fc=0):
          super(RNN_model, self).__init__()
          # embedding_dim: size of each embedding vector
          # hidden_size: number of features in the hidden state
          # num_layers: number of recurrent layers
          # bias: introduces a bias
          # batch_first: input and output tensors are provided as (batch, seq, feature)
          # dropout: if non-zero, introduces a dropout layer on the outputs of each RNN layer except the last
          ↪ layer

          self.embedding = nn.Embedding(num_embeddings=vocab_size, embedding_dim=embedding_dim,
          ↪ padding_idx=padding_idx)

          if type == "RNN":
              self.rnn = nn.RNN(input_size=embedding_dim, hidden_size=hidden_size, num_layers=num_layers,
          ↪ bias=True, batch_first=True, dropout=dropout_rnn, nonlinearity="tanh")
          elif type == "GRU":
              self.rnn = nn.GRU(input_size=embedding_dim, hidden_size=hidden_size, num_layers=num_layers,
          ↪ bias=True, batch_first=True, dropout=dropout_rnn)
          elif type == "LSTM":
              self.rnn = nn.LSTM(input_size=embedding_dim, hidden_size=hidden_size, num_layers=num_layers,
          ↪ bias=True, batch_first=True, dropout=dropout_rnn)

          self.fc = nn.Linear(in_features=hidden_size, out_features=num_classes)
          self.dropout = nn.Dropout(p=dropout_fc)

      def forward(self, x):
          x = self.embedding(x)
          x, _ = self.rnn(x)
```

```

x = x[:, -1, :] # extract last hidden state for each sequence
x = self.dropout(x) # apply dropout to the last hidden state
x = self.fc(x) # pass last hidden state through the fully connected layer
return x

```

2.2 Hyper parameter tuning

```

[73]: train_dataset = EmotionDataset(train_df["padded"].tolist(), train_df["label"].tolist())
      val_dataset = EmotionDataset(val_df["padded"].tolist(), val_df["label"].tolist())
      test_dataset = EmotionDataset(test_df["padded"].tolist(), test_df["label"].tolist())

```

```

[ ]: def grid_search(vocab_size, num_classes, train_dataset, val_dataset, label_map, device):
      param_grid = {
          # The best hyperparameters found is commented out
          "type": ["LSTM", "GRU", "RNN"], # "LSTM"
          "embedding_dim": [100, 75], # 75
          "hidden_size": [512, 256], # 256
          "layers": [1, 2], # 1
          "dropout_rnn": [0.0, 0.2], # 0.0
          "dropout_fc": [0.4, 0.6], # 0.4
          "learning_rate": [0.001, 0.0005], # 0.0001
          "reg_lambda": [0.0001, 0.00005], # 0.0001
          "batch_size": [16, 32], # 16
      }

      # Generate all combinations of hyperparameters
      keys, values = zip(*param_grid.items())
      configs = [dict(zip(keys, v)) for v in itertools.product(*values)]
      total_configs = len(configs)

      best_val_f1_score = 0
      results_list = []
      for i, config in enumerate(configs):
          label = f"model_{i}"
          print(f"Training model: {label} ({i+1}/{total_configs})")
          print(json.dumps(config, indent=4))

          train_loader = DataLoader(train_dataset, batch_size=config["batch_size"], shuffle=True)
          val_loader = DataLoader(val_dataset, batch_size=config["batch_size"], shuffle=False)

          # Model
          model = RNN_model(
              type=config["type"],
              vocab_size=vocab_size,
              embedding_dim=config["embedding_dim"],
              hidden_size=config["hidden_size"],
              num_classes=num_classes,
              num_layers=config["layers"],
              dropout_rnn=config["dropout_rnn"],
              dropout_fc=config["dropout_fc"],
          )

          # Train
          results = train_model(
              label=label,
              model=model,
              train_loader=train_loader,
              val_loader=val_loader,
              label_map=label_map,
              device=device,
              optimizer_type="AdamW",
              learning_rate=config["learning_rate"],
              reg_type="L2",
              reg_lambda=config["reg_lambda"],
              num_epochs=10
          )

```

```

        # Track the best model
        current_val_f1_score = max(results["val_f1_scores"])
        if current_val_f1_score > best_val_f1_score:
            best_val_f1_score = current_val_f1_score
            best_model_label = label
            print(f"New best model found: {best_model_label} with val f1 score: {best_val_f1_score:.4f}")

        # Add results to list
        results["config"] = config
        results_list.append(results)

    # Save results
    with open("results/grid_search_results.json", "w") as f:
        json.dump(results_list, f, indent=4)

    print(f"Best model found: {best_model_label}")

grid_search(vocab_size, num_classes, train_dataset, val_dataset, label_map, device)

```

2.3 Best RNN model

```

[75]: label = "best_rnn_model"
      best_batch_size = 16

      train_loader = DataLoader(train_dataset, batch_size=best_batch_size, shuffle=True)
      val_loader = DataLoader(val_dataset, batch_size=best_batch_size, shuffle=False)

      best_rnn_model = RNN_model(
          type="LSTM",
          vocab_size=vocab_size,
          embedding_dim=75,
          hidden_size=256,
          num_classes=num_classes,
          num_layers=1,
          dropout_rnn=0.0,
          dropout_fc=0.4,
      )

      best_rnn_results = train_model(
          label=label,
          model=best_rnn_model,
          train_loader=train_loader,
          val_loader=val_loader,
          label_map=label_map,
          device=device,
          optimizer_type="AdamW",
          learning_rate=0.001,
          reg_type="L2",
          reg_lambda=0.0001,
          num_epochs=4
      )

      plot_scores(best_rnn_results, label)

```

```

Epoch 1/4 (10s) | Train: loss 43.7021, acc 40.13%, f1 34.61% | Val: loss 1.1907,
acc 57.05%, f1 54.13%
Epoch 2/4 (7s) | Train: loss 13.1735, acc 69.10%, f1 68.43% | Val: loss 0.6395,
acc 80.35%, f1 80.08%
Epoch 3/4 (7s) | Train: loss 5.0673, acc 83.09%, f1 82.97% | Val: loss 0.5178,
acc 83.10%, f1 82.92%
Epoch 4/4 (7s) | Train: loss 2.5631, acc 86.90%, f1 86.83% | Val: loss 0.551,
acc 82.20%, f1 82.18%
Total Training Time: 32s

accuracy: 0.822
f1: 0.8218
precision: 0.8257

```

```

recall: 0.822
classification report:


```

	precision	recall	f1-score	support
sadness	0.80	0.91	0.85	550
joy	0.89	0.83	0.86	704
love	0.69	0.69	0.69	178
anger	0.86	0.76	0.80	275
fear	0.79	0.80	0.79	212
surprise	0.75	0.77	0.76	81
accuracy			0.82	2000
macro avg	0.79	0.79	0.79	2000
weighted avg	0.83	0.82	0.82	2000

```

confusion matrix:
[502, 15, 7, 8, 16, 2]
[54, 581, 42, 10, 10, 7]
[15, 29, 122, 6, 3, 3]
[41, 14, 3, 208, 8, 1]
[16, 6, 4, 9, 169, 8]
[3, 5, 0, 2, 9, 62]

```



3 Task 3: Transformer model

```

[76]: batch_size = 32
train_dataset = EmotionDataset(train_df["padded"].tolist(), train_df["label"].tolist())
val_dataset = EmotionDataset(val_df["padded"].tolist(), val_df["label"].tolist())
test_dataset = EmotionDataset(test_df["padded"].tolist(), test_df["label"].tolist())
train_labels = train_df["label"].tolist()
val_labels = val_df["label"].tolist()
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
n_embeds = 10336
num_epochs = 50

```

```

[81]: config = {
    "label": "transformer1",
    "d_key": 64,
    "n_heads": 8,
    "mlp_factor": 4,
    "d_model": 128,
    "n_layers": 6,
    "dropout1": 0.5,
    "dropout2": 0.25,
    "optimizer_type": "Adam",
    "learning_rate": 0.0003,
    "weight_decay": 0,

```

```

        "reg_type": "L1",
        "reg_lambda": 1e-4,
    }
    print(f"Training model: {config['label']}")
    model = transformer_model.TransformerClassifier(n_embeds=n_embeds, n_classes=6, d_model=config["d_model"],
    ↪d_key=config["d_key"], n_heads=config["n_heads"], mlp_factor=config["mlp_factor"],
    ↪n_layers=config["n_layers"], device = device, dropout1=config["dropout1"], dropout2=config["dropout2"])
    results = train_model(
        label=config["label"],
        model=model,
        train_loader=train_loader,
        val_loader=val_loader,
        label_map=label_map,
        device=device,
        optimizer_type=config["optimizer_type"],
        learning_rate=config["learning_rate"],
        weight_decay=config["weight_decay"],
        reg_type=config["reg_type"],
        reg_lambda=config["reg_lambda"],
        num_epochs=num_epochs,
    )
    plot_scores(results, config["label"])

```

```

Training model: transformer1
Epoch 1/50 (89s) | Train: loss 102.0006, acc 33.86%, f1 23.71% | Val: loss
1.5793, acc 32.10%, f1 23.49%
Epoch 2/50 (76s) | Train: loss 82.9869, acc 36.89%, f1 27.59% | Val: loss
1.5281, acc 40.85%, f1 31.31%
Epoch 3/50 (76s) | Train: loss 69.0236, acc 41.38%, f1 33.03% | Val: loss
1.4376, acc 45.90%, f1 38.03%
Epoch 4/50 (80s) | Train: loss 57.1575, acc 46.12%, f1 41.04% | Val: loss
1.3466, acc 50.80%, f1 46.88%
Epoch 5/50 (77s) | Train: loss 47.2545, acc 52.08%, f1 48.79% | Val: loss
1.1978, acc 54.85%, f1 50.44%
Epoch 6/50 (80s) | Train: loss 39.011, acc 59.49%, f1 57.57% | Val: loss 1.062,
acc 62.00%, f1 60.50%
Epoch 7/50 (79s) | Train: loss 32.3509, acc 65.00%, f1 63.95% | Val: loss
0.9193, acc 68.25%, f1 67.81%
Epoch 8/50 (77s) | Train: loss 26.9274, acc 70.02%, f1 69.34% | Val: loss
0.8478, acc 70.40%, f1 69.82%
Epoch 9/50 (76s) | Train: loss 22.5769, acc 73.58%, f1 73.13% | Val: loss
0.8064, acc 71.35%, f1 71.28%
Epoch 10/50 (76s) | Train: loss 19.0561, acc 76.75%, f1 76.42% | Val: loss
0.712, acc 74.90%, f1 74.65%
Epoch 11/50 (76s) | Train: loss 16.2046, acc 78.24%, f1 77.94% | Val: loss
0.6758, acc 77.20%, f1 77.05%
Epoch 12/50 (76s) | Train: loss 13.8662, acc 79.91%, f1 79.69% | Val: loss
0.6526, acc 77.35%, f1 76.76%
Epoch 13/50 (76s) | Train: loss 11.9736, acc 81.45%, f1 81.25% | Val: loss
0.6205, acc 79.95%, f1 79.66%
Epoch 14/50 (77s) | Train: loss 10.4343, acc 82.03%, f1 81.87% | Val: loss
0.6295, acc 78.65%, f1 78.57%
Epoch 15/50 (77s) | Train: loss 9.1926, acc 82.99%, f1 82.84% | Val: loss
0.5959, acc 79.95%, f1 79.67%
Epoch 16/50 (76s) | Train: loss 8.1612, acc 83.69%, f1 83.55% | Val: loss
0.5867, acc 80.30%, f1 80.02%
Epoch 17/50 (77s) | Train: loss 7.3331, acc 83.83%, f1 83.73% | Val: loss
0.5698, acc 81.00%, f1 80.82%
Epoch 18/50 (76s) | Train: loss 6.6518, acc 84.61%, f1 84.49% | Val: loss
0.6136, acc 79.55%, f1 79.36%
Epoch 19/50 (76s) | Train: loss 6.104, acc 84.82%, f1 84.73% | Val: loss 0.5767,
acc 80.55%, f1 80.34%
Epoch 20/50 (76s) | Train: loss 5.6224, acc 85.53%, f1 85.45% | Val: loss
0.5924, acc 80.15%, f1 79.88%
Epoch 21/50 (76s) | Train: loss 5.2473, acc 85.56%, f1 85.48% | Val: loss
0.6482, acc 78.85%, f1 78.85%
Epoch 22/50 (76s) | Train: loss 4.9044, acc 86.11%, f1 86.02% | Val: loss
0.5677, acc 80.80%, f1 80.52%

```

Epoch 23/50 (76s) | Train: loss 4.6268, acc 86.14%, f1 86.07% | Val: loss 0.5594, acc 82.25%, f1 82.19%

Epoch 24/50 (76s) | Train: loss 4.3677, acc 86.46%, f1 86.37% | Val: loss 0.5706, acc 81.20%, f1 80.95%

Epoch 25/50 (78s) | Train: loss 4.1486, acc 86.97%, f1 86.90% | Val: loss 0.5459, acc 81.90%, f1 81.75%

Epoch 26/50 (76s) | Train: loss 3.9806, acc 86.69%, f1 86.62% | Val: loss 0.542, acc 82.40%, f1 82.38%

Epoch 27/50 (77s) | Train: loss 3.7892, acc 87.52%, f1 87.46% | Val: loss 0.542, acc 82.30%, f1 81.97%

Epoch 28/50 (76s) | Train: loss 3.6444, acc 87.35%, f1 87.29% | Val: loss 0.5416, acc 82.55%, f1 82.39%

Epoch 29/50 (76s) | Train: loss 3.4955, acc 87.80%, f1 87.74% | Val: loss 0.5703, acc 81.15%, f1 80.88%

Epoch 30/50 (77s) | Train: loss 3.389, acc 87.71%, f1 87.65% | Val: loss 0.546, acc 82.00%, f1 81.85%

Epoch 31/50 (76s) | Train: loss 3.2559, acc 88.16%, f1 88.10% | Val: loss 0.5646, acc 82.30%, f1 81.96%

Epoch 32/50 (76s) | Train: loss 3.1771, acc 87.64%, f1 87.59% | Val: loss 0.5844, acc 81.15%, f1 81.21%

Epoch 33/50 (76s) | Train: loss 3.0501, acc 88.59%, f1 88.54% | Val: loss 0.5427, acc 82.45%, f1 82.40%

Epoch 34/50 (76s) | Train: loss 2.9593, acc 88.60%, f1 88.55% | Val: loss 0.5736, acc 81.90%, f1 81.82%

Epoch 35/50 (76s) | Train: loss 2.8801, acc 88.61%, f1 88.56% | Val: loss 0.5598, acc 82.25%, f1 82.16%

Epoch 36/50 (77s) | Train: loss 2.8071, acc 88.83%, f1 88.78% | Val: loss 0.5595, acc 82.15%, f1 82.04%

Epoch 37/50 (77s) | Train: loss 2.7384, acc 88.84%, f1 88.79% | Val: loss 0.6097, acc 81.05%, f1 80.46%

Epoch 38/50 (76s) | Train: loss 2.659, acc 89.01%, f1 88.96% | Val: loss 0.5831, acc 81.10%, f1 80.82%

Epoch 39/50 (77s) | Train: loss 2.5879, acc 88.94%, f1 88.89% | Val: loss 0.5631, acc 82.45%, f1 82.42%

Epoch 40/50 (77s) | Train: loss 2.5379, acc 89.44%, f1 89.40% | Val: loss 0.589, acc 82.45%, f1 82.54%

Epoch 41/50 (77s) | Train: loss 2.4701, acc 89.35%, f1 89.31% | Val: loss 0.5726, acc 82.80%, f1 82.71%

Epoch 42/50 (76s) | Train: loss 2.4347, acc 89.32%, f1 89.28% | Val: loss 0.5533, acc 83.55%, f1 83.54%

Epoch 43/50 (77s) | Train: loss 2.3524, acc 89.97%, f1 89.93% | Val: loss 0.5814, acc 82.50%, f1 82.40%

Epoch 44/50 (76s) | Train: loss 2.3024, acc 90.38%, f1 90.34% | Val: loss 0.5854, acc 83.15%, f1 83.07%

Epoch 45/50 (77s) | Train: loss 2.2664, acc 89.99%, f1 89.95% | Val: loss 0.5644, acc 83.30%, f1 83.14%

Epoch 46/50 (76s) | Train: loss 2.2204, acc 90.38%, f1 90.34% | Val: loss 0.5628, acc 82.60%, f1 82.43%

Epoch 47/50 (76s) | Train: loss 2.167, acc 90.64%, f1 90.61% | Val: loss 0.606, acc 81.65%, f1 81.72%

Epoch 48/50 (76s) | Train: loss 2.1392, acc 90.64%, f1 90.62% | Val: loss 0.5831, acc 83.00%, f1 82.94%

Epoch 49/50 (76s) | Train: loss 2.0786, acc 91.10%, f1 91.07% | Val: loss 0.5902, acc 82.90%, f1 82.65%

Epoch 50/50 (76s) | Train: loss 2.056, acc 90.95%, f1 90.92% | Val: loss 0.594, acc 82.80%, f1 82.62%

Total Training Time: 3843s

accuracy: 0.828

f1: 0.8262

precision: 0.8267

recall: 0.828

classification report:

	precision	recall	f1-score	support
sadness	0.88	0.85	0.87	550
joy	0.83	0.89	0.86	704
love	0.74	0.60	0.66	178


anger	0.81	0.82	0.81	275
fear	0.79	0.81	0.80	212
surprise	0.74	0.68	0.71	81
accuracy			0.83	2000
macro avg	0.80	0.78	0.79	2000
weighted avg	0.83	0.83	0.83	2000

confusion matrix:

```

[470, 34, 7, 18, 18, 3]
[23, 628, 24, 17, 5, 7]
[8, 51, 107, 8, 3, 1]
[16, 17, 4, 225, 12, 1]
[14, 12, 2, 6, 171, 7]
[2, 12, 0, 4, 8, 55]

```



Train multiple:

```
[ ]: configurations = [
    #heads 16
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 16,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L2",
        "reg_lambda": 1e-4,
    },
    # mlp_factor 2
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 2,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L2",
    }
]
```

```

        "reg_lambda": 1e-4,
    },
    # 0,0003 learning rate
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.0003,
        "weight_decay": 0,
        "reg_type": "L2",
        "reg_lambda": 1e-4,
    },
    # L1 reg
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.5,
        "dropout2": 0.25,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L1",
        "reg_lambda": 1e-4,
    },
    # More dropout
    {
        "label": "transformer1",
        "d_key": 64,
        "n_heads": 8,
        "mlp_factor": 4,
        "d_model": 128,
        "n_layers": 6,
        "dropout1": 0.6,
        "dropout2": 0.3,
        "optimizer_type": "Adam",
        "learning_rate": 0.001,
        "weight_decay": 0,
        "reg_type": "L2",
        "reg_lambda": 1e-4,
    },
]
for config in configurations:
    print(f"Training model: {config['label']}")
    model = transformer_model.TransformerClassifier(n_embeds=n_embeds, n_classes=6,
    ↪d_model=config["d_model"], d_key=config["d_key"], n_heads=config["n_heads"],
    ↪mlp_factor=config["mlp_factor"], n_layers=config["n_layers"], device = device,
    ↪dropout1=config["dropout1"], dropout2=config["dropout2"])
    results = train_model(
        label=config["label"],
        model=model,
        train_loader=train_loader,
        val_loader=val_loader,
        label_map=label_map,
        device=device,
        optimizer_type=config["optimizer_type"],
        learning_rate=config["learning_rate"],
        weight_decay=config["weight_decay"],

```

```

        reg_type=config["reg_type"],
        reg_lambda=config["reg_lambda"],
        num_epochs=num_epochs,
    )
    plot_scores(results, config["label"])

```

4 Task 4: Analysis

```

[ ]: test_loader = DataLoader(test_dataset, batch_size=best_batch_size, shuffle=True)

predict("best_rnn_model", best_rnn_model, device, test_loader, label_map, reverse_vocab)

```

accuracy: 0.8245

f1: 0.8248

precision: 0.8261

recall: 0.8245

classification report:

	precision	recall	f1-score	support
sadness	0.85	0.90	0.87	581
joy	0.87	0.82	0.84	695
love	0.64	0.70	0.67	159
anger	0.83	0.81	0.82	275
fear	0.81	0.81	0.81	224
surprise	0.64	0.62	0.63	66
accuracy			0.82	2000
macro avg	0.77	0.78	0.78	2000
weighted avg	0.83	0.82	0.82	2000

confusion matrix:

```

[520, 31, 4, 12, 12, 2]
[35, 571, 52, 20, 8, 9]
[16, 23, 112, 4, 2, 2]
[20, 20, 4, 223, 7, 1]
[17, 6, 2, 8, 182, 9]
[4, 7, 0, 1, 13, 41]

```

```

[82]: test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True)

predict("best_rnn_model", model, device, test_loader, label_map, reverse_vocab)

```

accuracy: 0.813

f1: 0.8115

precision: 0.8126

recall: 0.813

classification report:

	precision	recall	f1-score	support
sadness	0.89	0.82	0.85	581
joy	0.81	0.88	0.85	695
love	0.68	0.59	0.63	159
anger	0.77	0.81	0.79	275
fear	0.80	0.80	0.80	224
surprise	0.67	0.56	0.61	66
accuracy			0.81	2000
macro avg	0.77	0.75	0.76	2000
weighted avg	0.81	0.81	0.81	2000

confusion matrix:

```

[479, 52, 6, 28, 13, 3]
[21, 614, 29, 15, 9, 7]
[7, 46, 94, 11, 0, 1]
[18, 19, 7, 222, 9, 0]
[15, 10, 2, 10, 180, 7]

```

```
[0, 13, 0, 3, 13, 37]
```

```
[ ]: config = {
    "label": "transformer1",
    "d_key": 64,
    "n_heads": 8,
    "mlp_factor": 4,
    "d_model": 128,
    "n_layers": 6,
    "dropout1": 0.5,
    "dropout2": 0.25,
    "optimizer_type": "Adam",
    "learning_rate": 0.0003,
    "weight_decay": 0,
    "reg_type": "L1",
    "reg_lambda": 1e-4,
}
model = transformer_model.TransformerClassifier(n_embeddings=n_embeddings, n_classes=6, d_model=config["d_model"],
↪ d_key=config["d_key"], n_heads=config["n_heads"], mlp_factor=config["mlp_factor"],
↪ n_layers=config["n_layers"], device = device, dropout1=config["dropout1"], dropout2=config["dropout2"])
model.load_state_dict(torch.load('models/transformer1.pth'))
predict_on_fly(model, tokenizer, vocab, device, label_map, max_length)
```

5 Task 5: Pre-trained model (transfer learning)

```
[ ]: model_name = "distilbert-base-uncased"
num_epochs = 5
learning_rate = 5e-6
max_length = 32 # different max length as we are using raw data to be tokenized using the pretrained
↪ tokenizer
batch_size = 32
label2id = {v: k for k, v in label_map.items()}

# Metric function
def pretrained_evaluation(predictions):
    preds = predictions.argmax(-1) # get predicted labels
    labels = predictions.label_ids
    accuracy = accuracy_score(labels, preds)
    f1 = f1_score(labels, preds, average="weighted")
    return {"accuracy": accuracy, "f1": f1}

# Load tokenizer
tokenizer_pretrained = AutoTokenizer.from_pretrained(model_name)

# Preprocessing function
def tokenize(batch):
    return tokenizer_pretrained(batch["text"], padding="max_length", truncation=True, max_length=max_length)

# Convert pandas DataFrame to Hugging Face Dataset
train_dataset = Dataset.from_pandas(train_df)
val_dataset = Dataset.from_pandas(val_df)
test_dataset = Dataset.from_pandas(test_df)

# Tokenize datasets
train_dataset = train_dataset.map(tokenize, batched=True)
val_dataset = val_dataset.map(tokenize, batched=True)
test_dataset = test_dataset.map(tokenize, batched=True)

# Set format for PyTorch
train_dataset = train_dataset.rename_column("label", "labels")
val_dataset = val_dataset.rename_column("label", "labels")
test_dataset = test_dataset.rename_column("label", "labels")
train_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
val_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
test_dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "labels"])
```

```

# Load pre-trained model
model = AutoModelForSequenceClassification.from_pretrained(
    model_name,
    num_labels=num_classes,
    id2label=label_map,
    label2id=label2id
)

# Training arguments
training_args = TrainingArguments(
    output_dir="results",
    eval_strategy="epoch", # evaluate at the end of each epoch
    save_strategy="epoch", # save checkpoints after every epoch
    logging_strategy="epoch",
    learning_rate=learning_rate,
    per_device_train_batch_size=batch_size,
    per_device_eval_batch_size=64,
    weight_decay=0.01,
    num_train_epochs=num_epochs,
    load_best_model_at_end=True, # load best model at the end of training
    metric_for_best_model="f1", # specify metric to monitor
    save_total_limit=1, # keep only the best checkpoint
)

# Define Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    compute_metrics=pretrained_evaluation, # compute metrics during evaluation
)

# Train model
trainer.train()

# Evaluate best model
best_results = trainer.evaluate()
print(f"Best Model Evaluation Results:\n {best_results}")

# Save best fine-tuned model
fine_tuned_model_name = "distilbert_finetuned"
trainer.save_model(f"models/{fine_tuned_model_name}")
tokenizer_pretrained.save_pretrained(f"models/{fine_tuned_model_name}")

```

```

[32]: # Load fine-tuned model
tokenizer_new = AutoTokenizer.from_pretrained(f"models/{fine_tuned_model_name}")
model_new = AutoModelForSequenceClassification.from_pretrained(f"models/{fine_tuned_model_name}",
    ↪ num_labels=num_classes)
trainer = Trainer(model=model_new)

# Predict on test set
predictions = trainer.predict(test_dataset)

# Extract logits and compute predicted labels
logits = torch.tensor(predictions.predictions) # convert logits to a PyTorch tensor
predicted_labels = torch.argmax(logits, dim=1).numpy() # convert to numpy array for sklearn metrics

# Evaluate predictions
metrics = compute_metrics(test_df["label"], predicted_labels, label_map.values())
print_metrics(metrics)
save_metrics(fine_tuned_model_name, metrics)

```

100%| 250/250 [00:01<00:00, 152.01it/s]

accuracy: 0.8885

f1: 0.8871

```

precision: 0.8869
recall: 0.8885
classification report:
      precision    recall  f1-score   support

    sadness      0.92      0.93      0.93      581
       joy      0.89      0.92      0.91      695
       love      0.74      0.70      0.72      159
       anger      0.90      0.90      0.90      275
        fear      0.90      0.87      0.88      224
    surprise      0.80      0.61      0.69       66

 accuracy          0.89      2000
  macro avg      0.86      0.82      0.84      2000
weighted avg      0.89      0.89      0.89      2000

confusion matrix:
[543, 20, 3, 11, 4, 0]
[14, 641, 32, 6, 1, 1]
[6, 39, 111, 3, 0, 0]
[14, 7, 2, 247, 5, 0]
[11, 2, 0, 7, 195, 9]
[4, 9, 1, 0, 12, 40]

```